NVIDIA NVIDIA DEEP LEARNING INSTITUTE

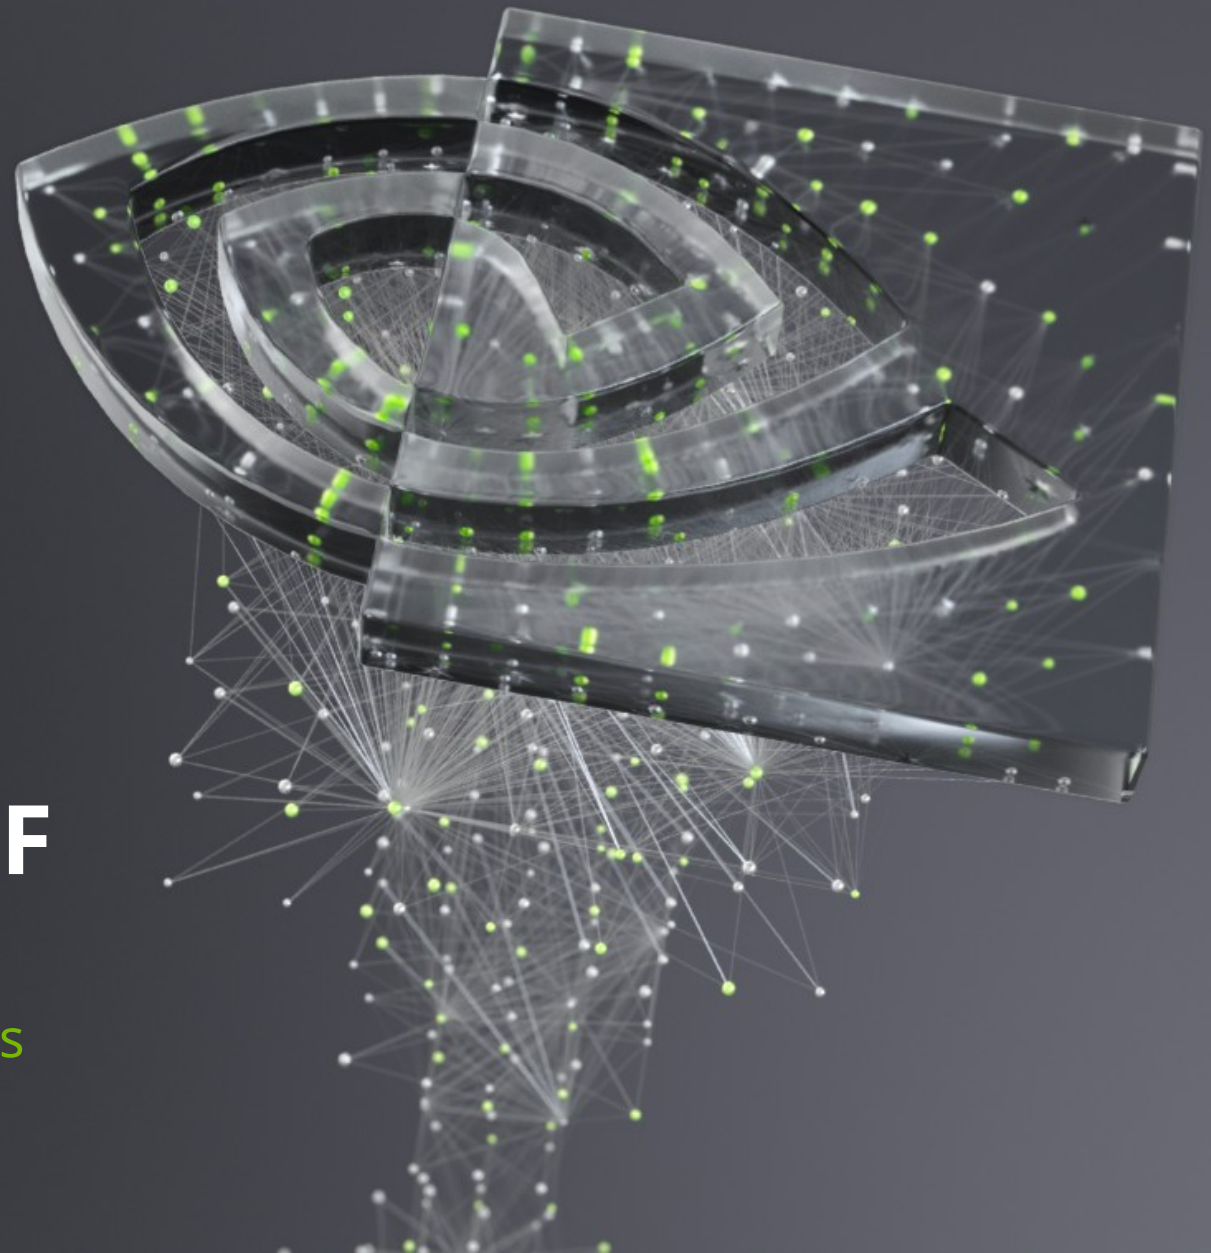

FUNDAMENTALS OF DEEP LEARNING

Part 2: How a Neural Network Trains

AGENDA

- Part 1: An Introduction to Deep Learning
- Part 2: How a Neural Network Trains
- Part 3: Convolutional Neural Networks
- Part 4: Data Augmentation and Deployment
- Part 5: Pre-trained Models
- Part 6: Advanced Architectures

AGENDA – PART 2

- Recap
- A Simpler Model
- From Neuron to Network
- Activation Functions
- Overfitting
- From Neuron to Classification

RECAP OF THE EXERCISE

What just happened?

- Loaded and visualized our data
- Edited our data (reshaped, normalized, to categorical)
- Created our model
- Compiled our model
- Trained the model on our data

DATA PREPARATION

Input as an array

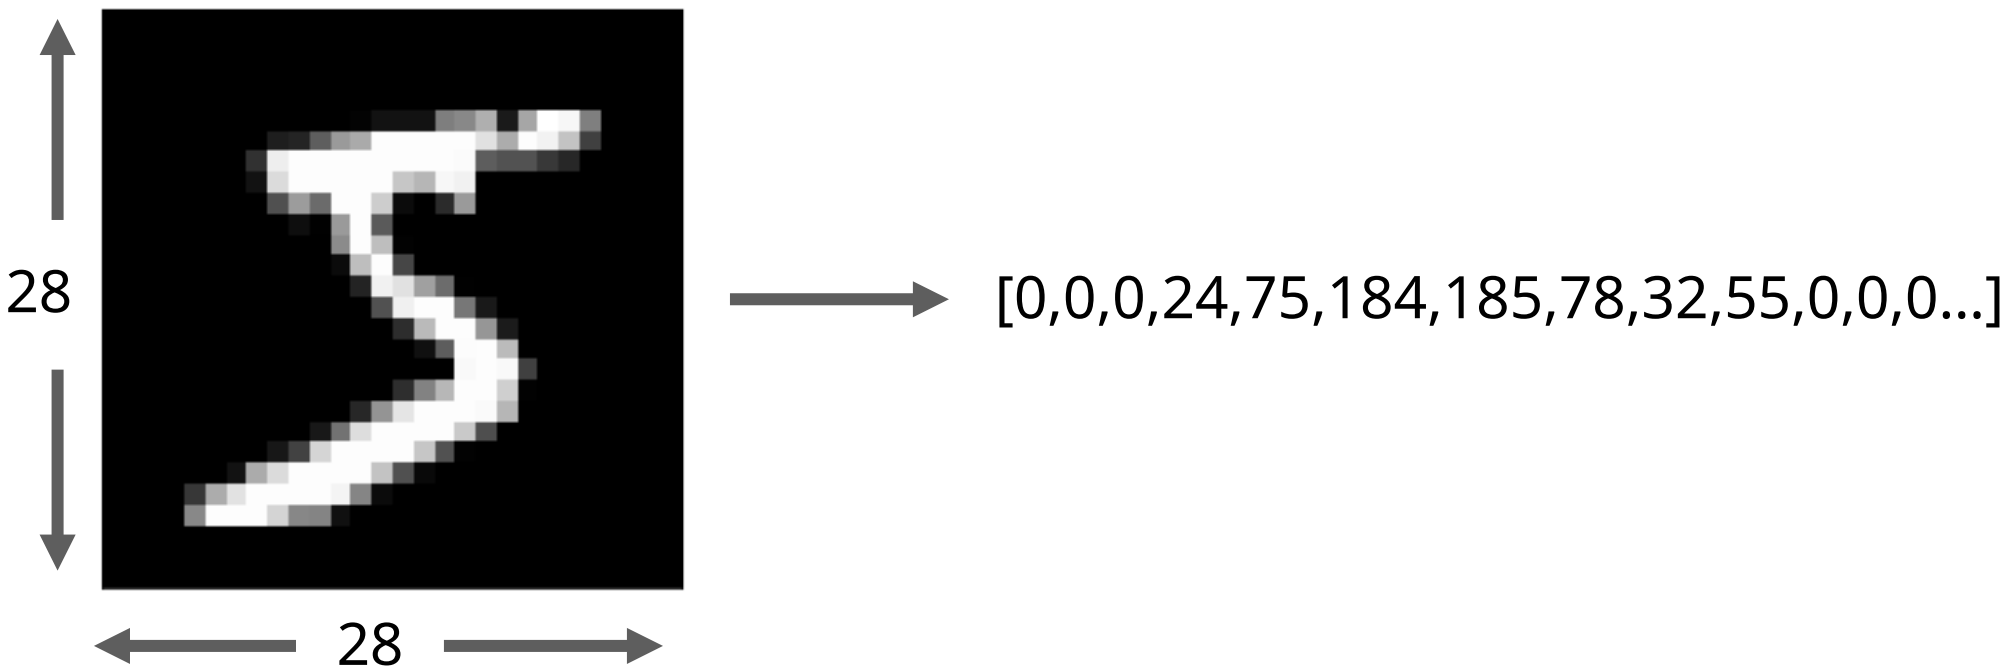

DATA PREPARATION

Targets as categories

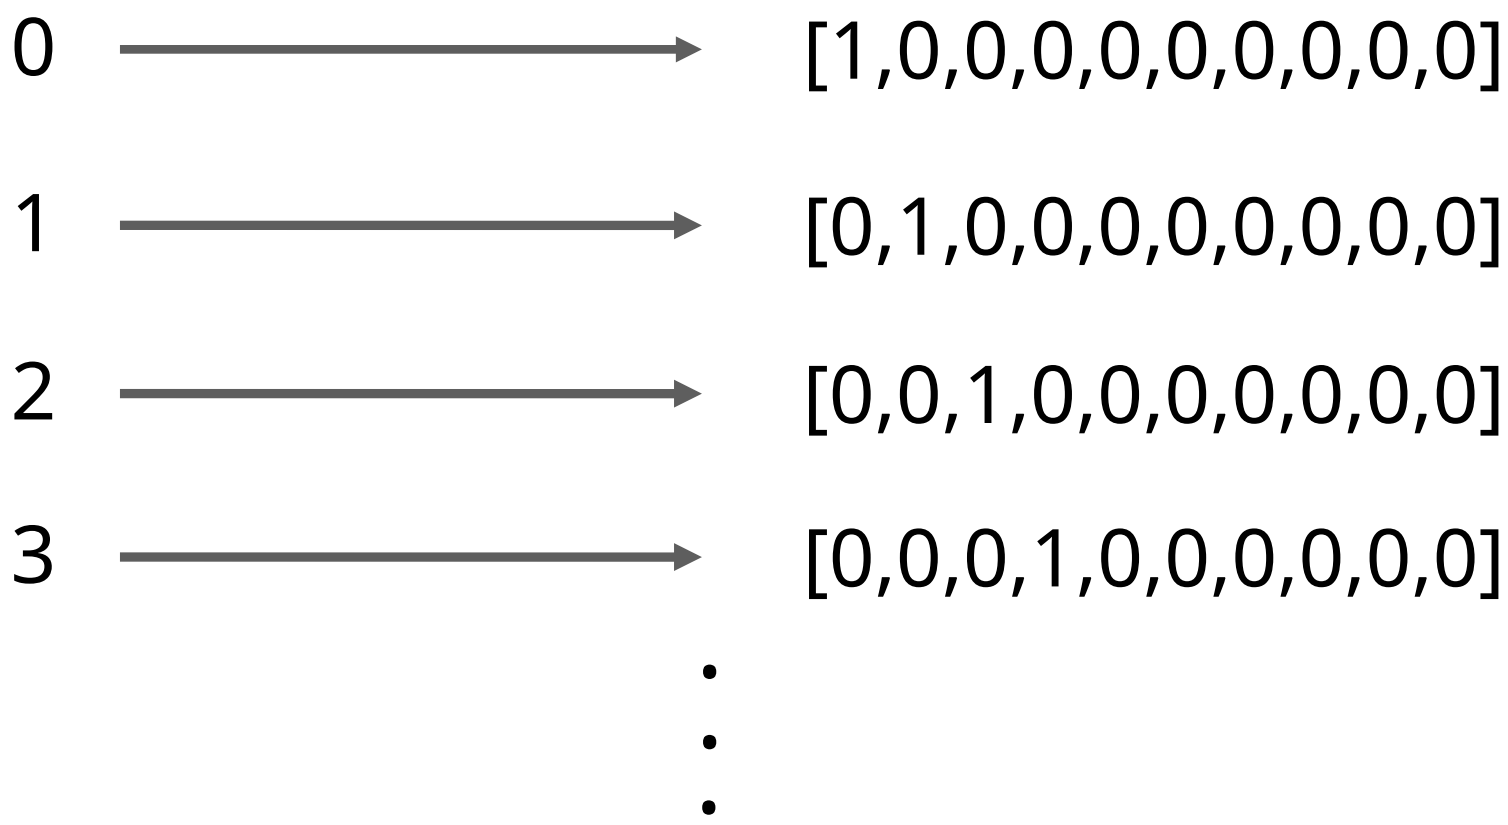

AN UNTRAINED MODEL

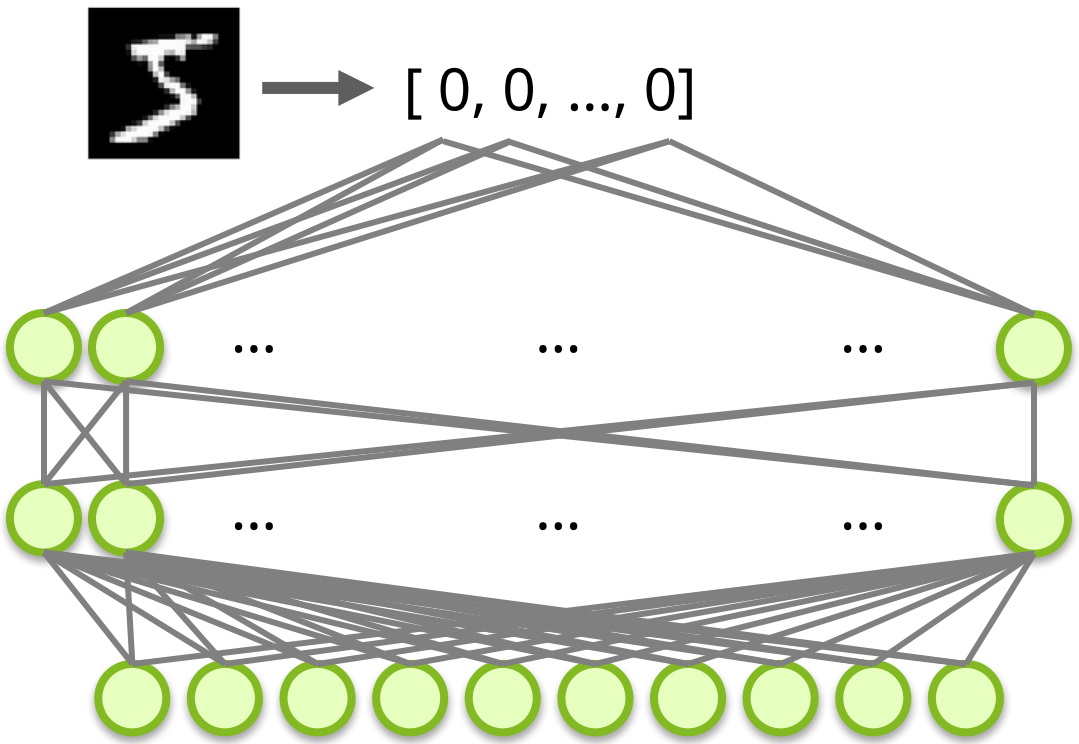

(784,)

(512,)

(512,)

(10,)

Layer Size

NVIDIA DEEP LEARNING INSTITUTE

A SIMPLER MODEL

A SIMPLER MODEL

$$y = mx + b$$

x	y
1	3
2	5

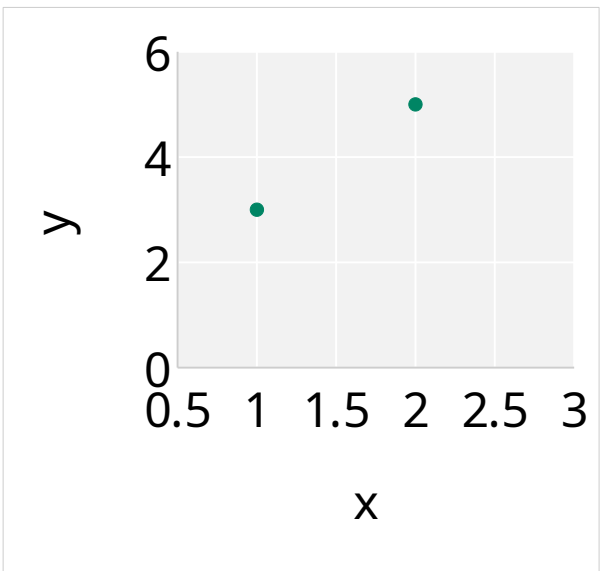

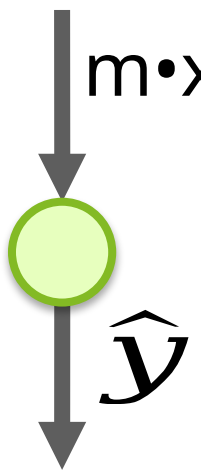

$m = ?$

$b = ?$

NVIDIA DEEP LEARNING INSTITUTE

A SIMPLER MODEL

$$y = mx + b$$

x	y
1	3
2	5

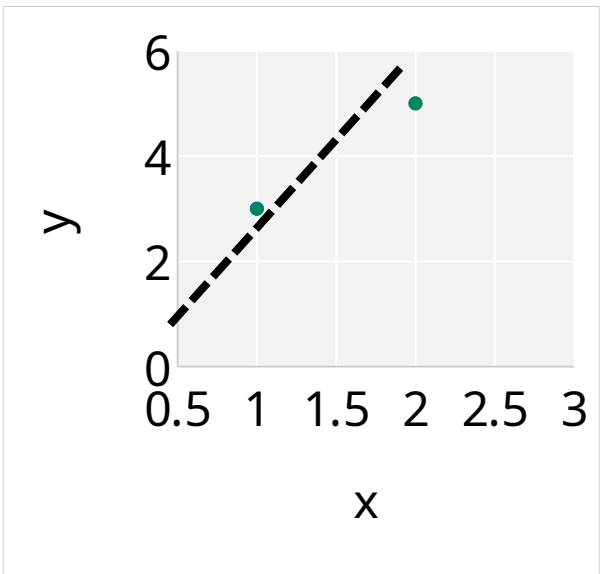

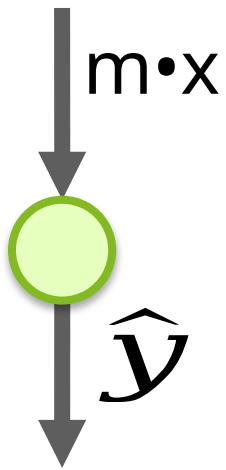

$m = ?$

$b = ?$

NVIDIA DEEP LEARNING INSTITUTE

A SIMPLER MODEL

$$y = mx + b$$

x	y	
1	3	4
2	5	3

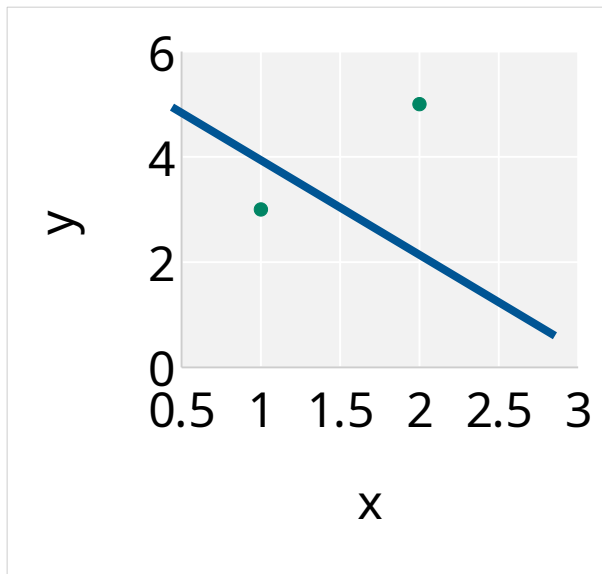

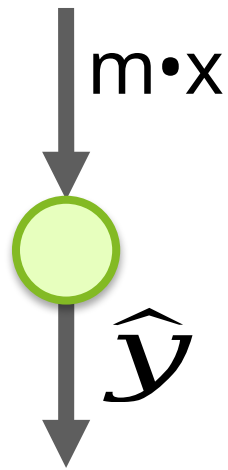

Start Random

$$m = -1$$

$$b = 5$$

NVIDIA DEEP LEARNING INSTITUTE

A SIMPLER MODEL

$$y = mx + b$$

x	y		
1	3	4	1
2	5	3	4
		MSE =	2.5
		RMSE =	1.6

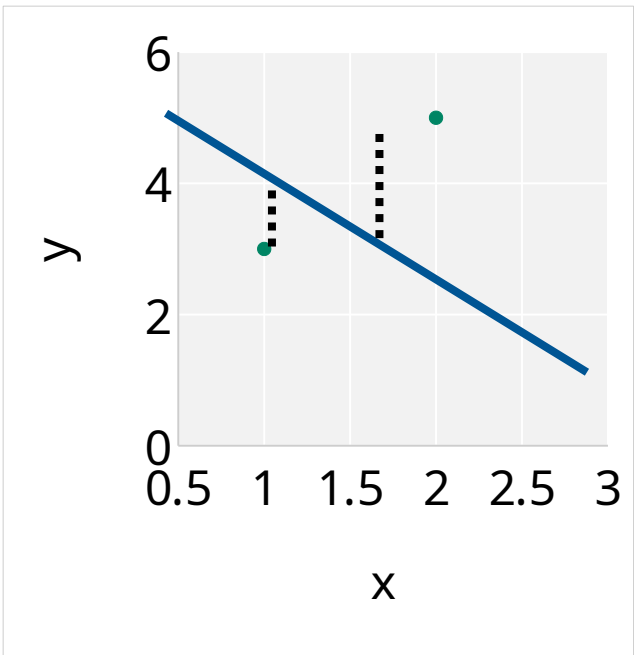

$$RMSE = \sqrt{\frac{1}{n}\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}$$

NVIDIA DEEP LEARNING INSTITUTE

A SIMPLER MODEL

$$y = mx + b$$

x	y		
1	3	4	1
2	5	3	4
		MSE =	2.5
		RMSE =	1.6

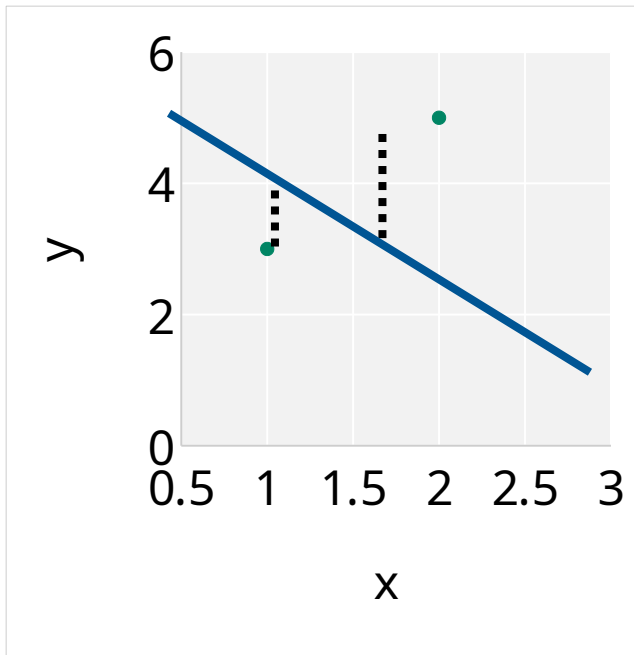


```python
data = [(1, 3), (2, 5)]
m = -1
b = 5


def get_rmse(data, m, b):
    """Calculates Mean Square Error"""
    n = len(data)
    squared_error = 0
    for x, y in data:
        # Find predicted y
        y_hat = m*x+b
        # Square difference between
        # prediction and true value
        squared_error += (
            y - y_hat)**2
    # Get average squared difference
    mse = squared_error / n
    # Square root for original units
    return mse ** .5

```


NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

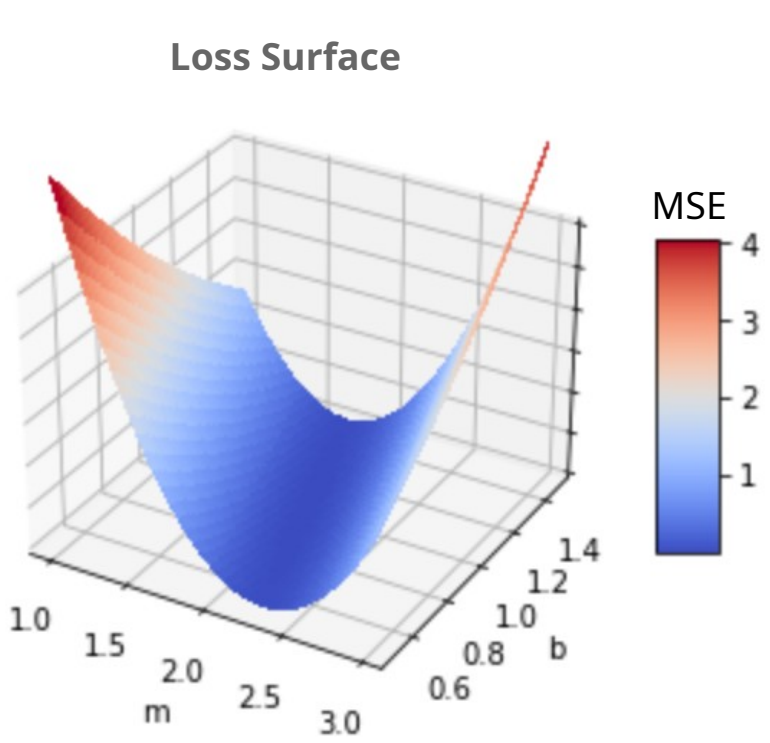

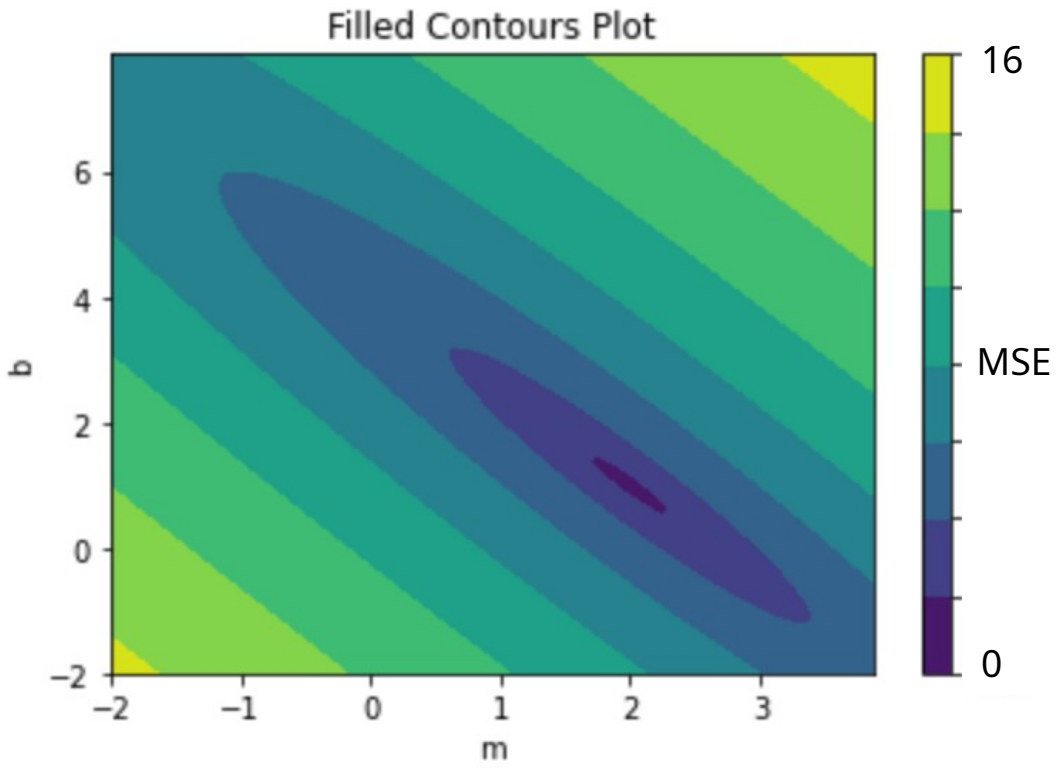

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

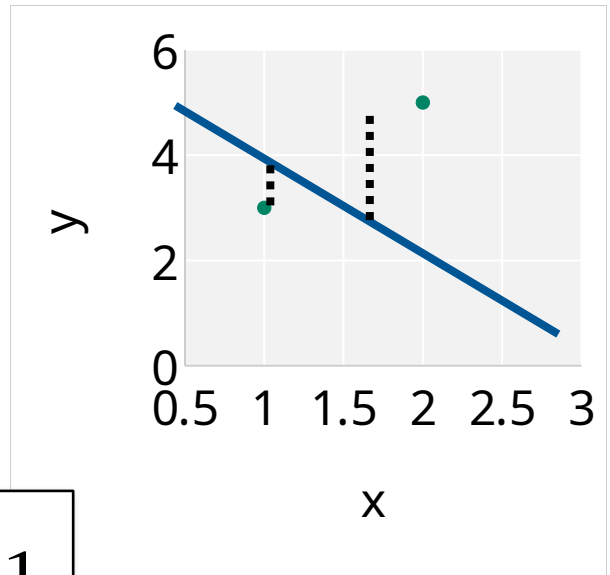

$$m = -1$$

$$b = 5$$

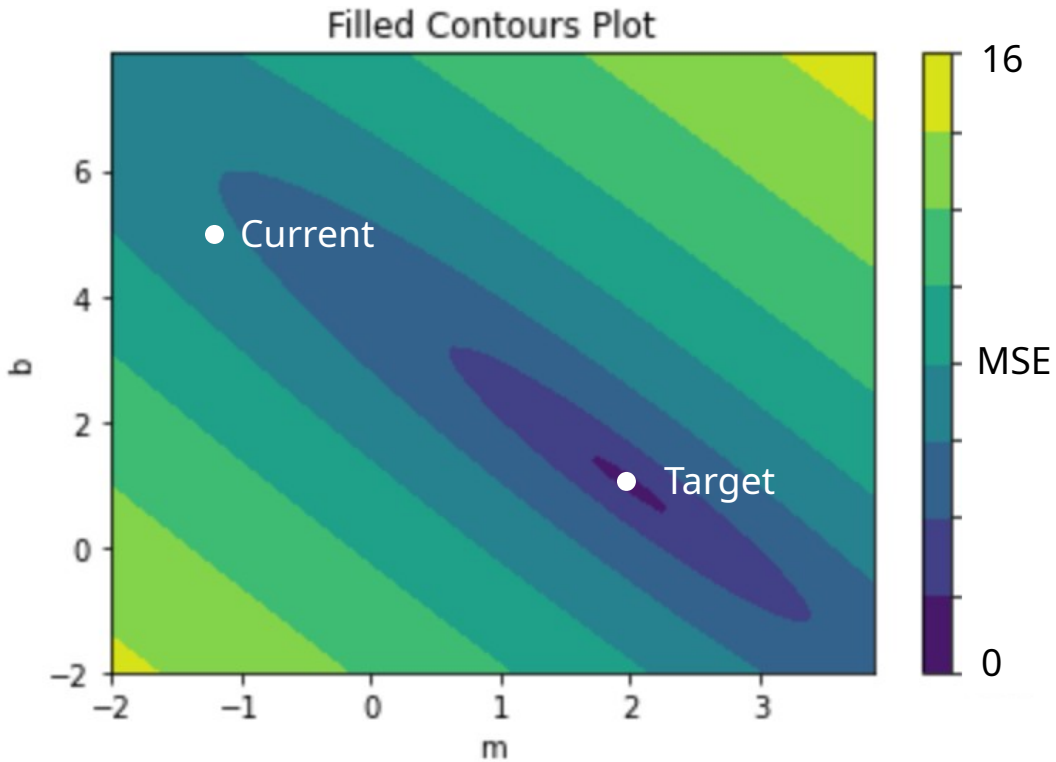

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

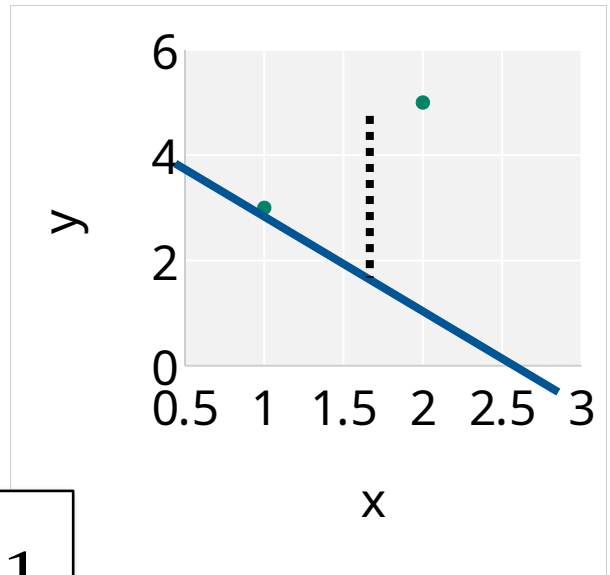

$$m = -1$$

$$b = 4$$

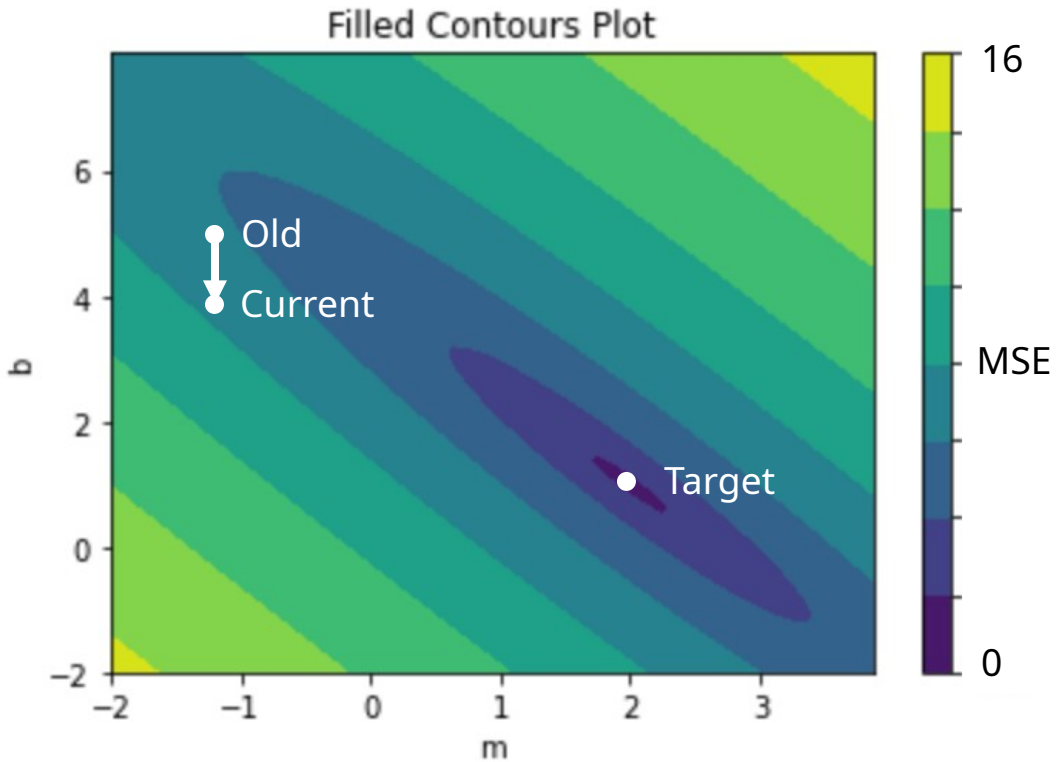

THE LOSS CURVE

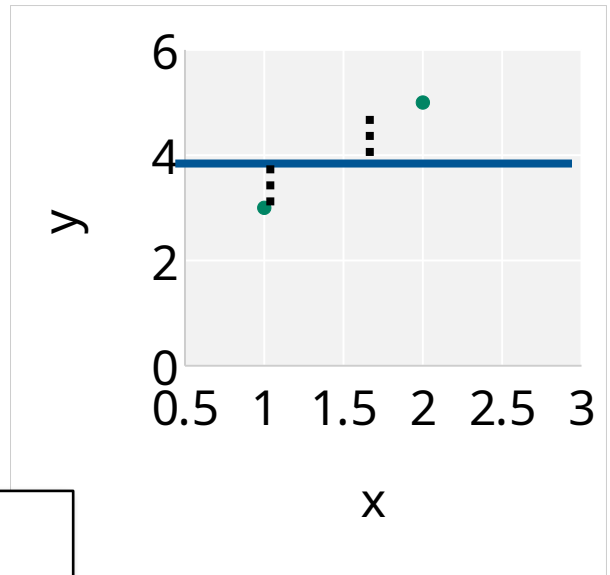

$$m = 0$$

$$b = 4$$

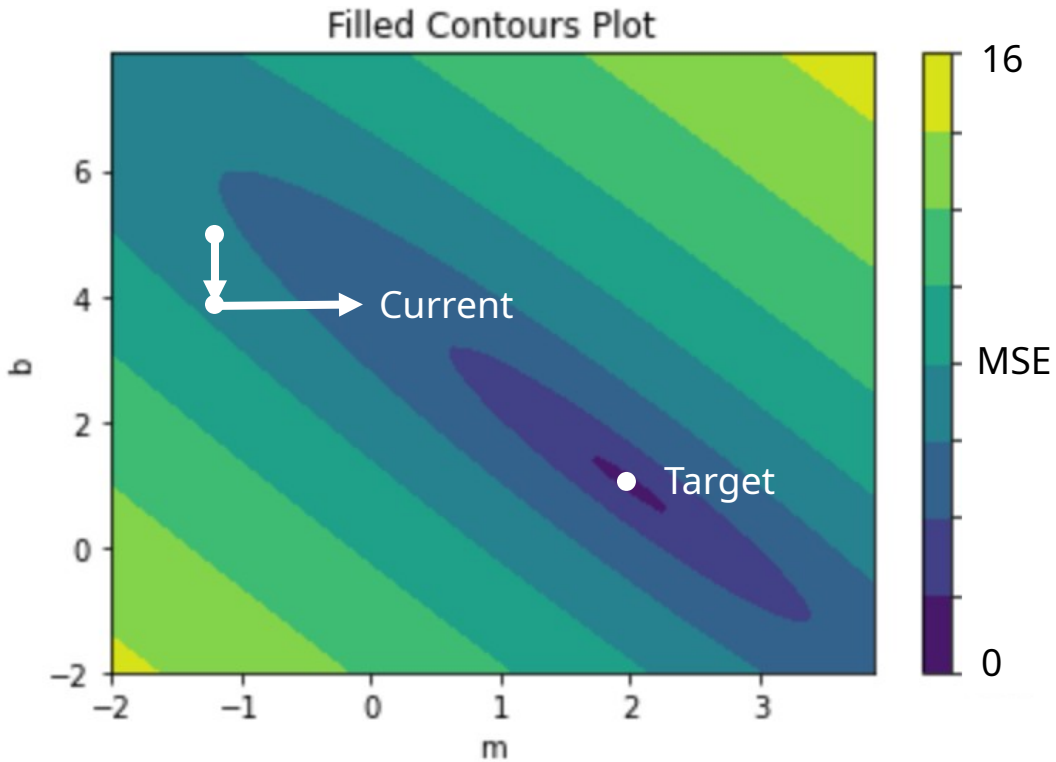

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

The Gradient	Which direction loss decreases the most
: The learning rate	How far to travel
Epoch	A model update with the full dataset
Batch	A sample of the full dataset
Step	An update to the weight parameters

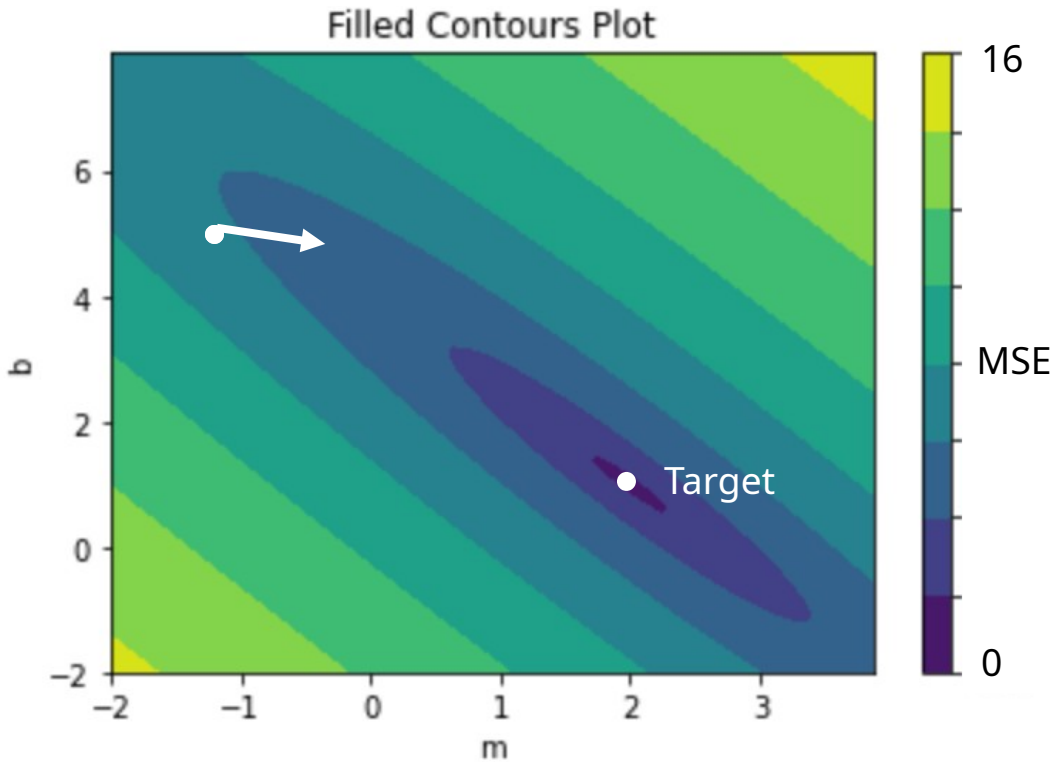

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

The Gradient	Which direction loss decreases the most
: The learning rate	How far to travel
Epoch	A model update with the full dataset
Batch	A sample of the full dataset
Step	An update to the weight parameters

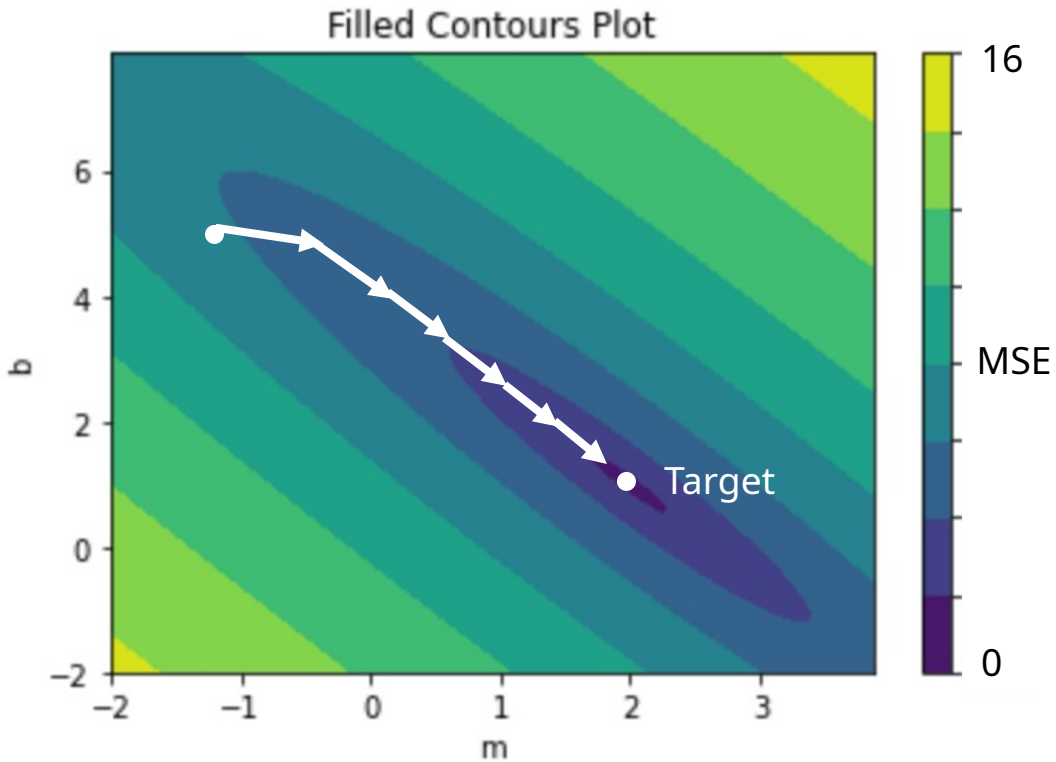

OPTIMIZERS

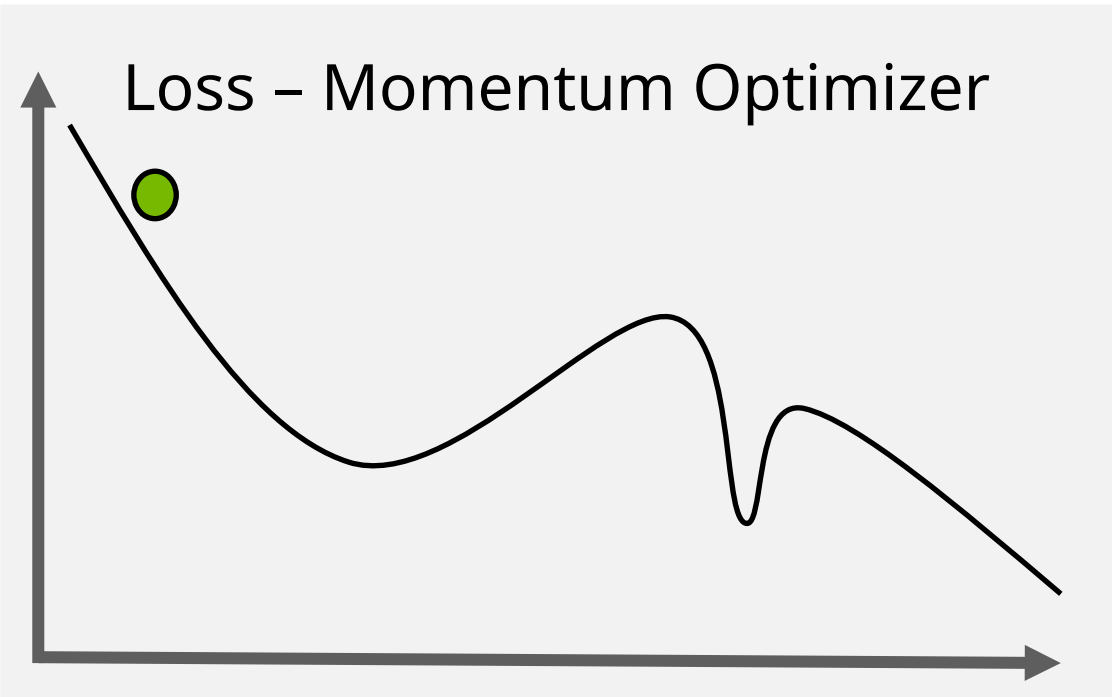

- Adam
- Adagrad
- RMSprop
- SGD

NVIDIA DEEP LEARNING INSTITUTE

FROM NEURON TO NETWORK

BUILDING A NETWORK

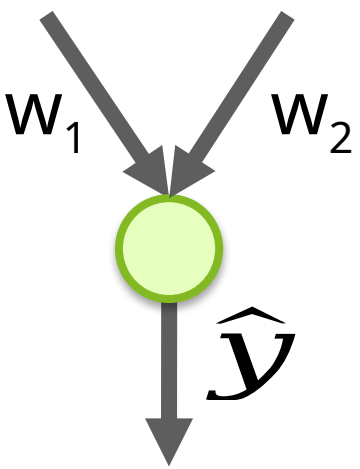

- Scales to more inputs

NVIDIA DEEP LEARNING INSTITUTE

BUILDING A NETWORK

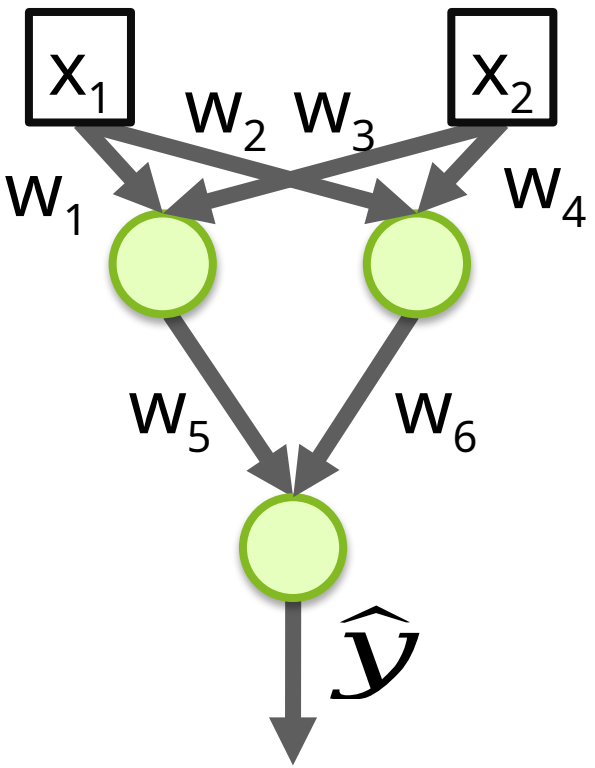

- Scales to more inputs
- Can chain neurons

NVIDIA DEEP LEARNING INSTITUTE

BUILDING A NETWORK

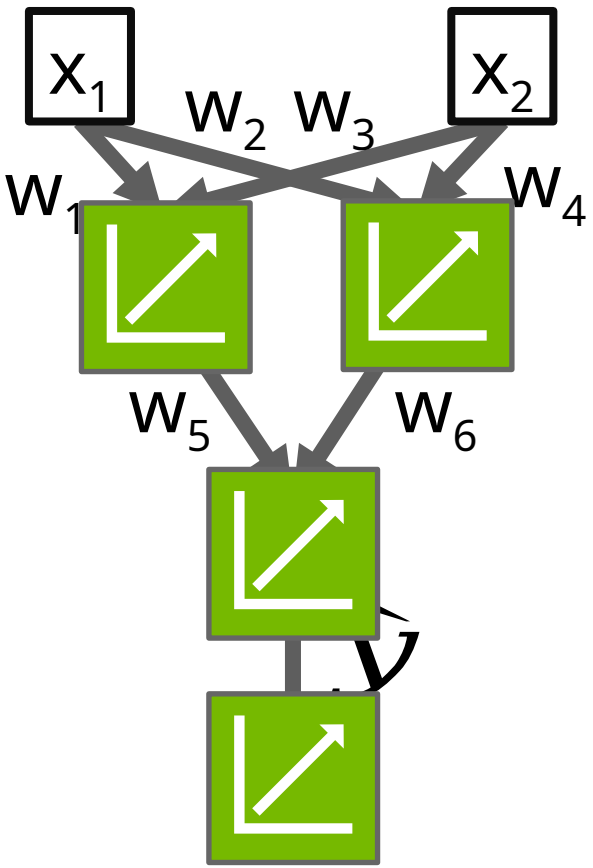

- Scales to more inputs
- Can chain neurons
- If all regressions are linear, then output will also be a linear regression

NVIDIA DEEP LEARNING INSTITUTE

ACTIVATION FUNCTIONS

ACTIVATION FUNCTIONS

Linear

$$\hat{y} = wx + b$$

```
# Multiply each input
# with a weight (w) and
# add intercept (b)
y_hat = wx+b
```

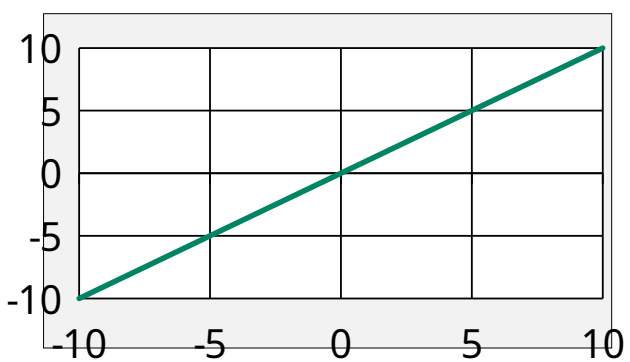

ReLU

$$\hat{y} = \begin{cases} wx+b \; if \; wx+b>0 \\ 0 \; otherwise \end{cases}$$

```
# Only return result
# if total is positive
linear = wx+b
y_hat = linear * (linear > 0)
```

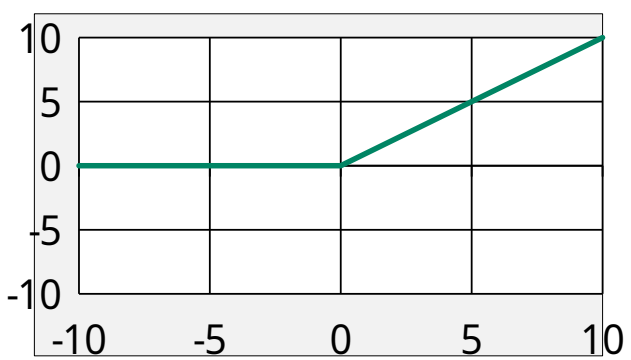

Sigmoid

$$\hat{y} = \frac{1}{1+e^{-(wx+b)}}$$

```
# Start with line
linear = wx + b
# Warp to - inf to 0
inf_to_zero = np.exp(-1 * linear)
# Squish to -1 to 1
y_hat = 1 / (1 + inf_to_zero)
```

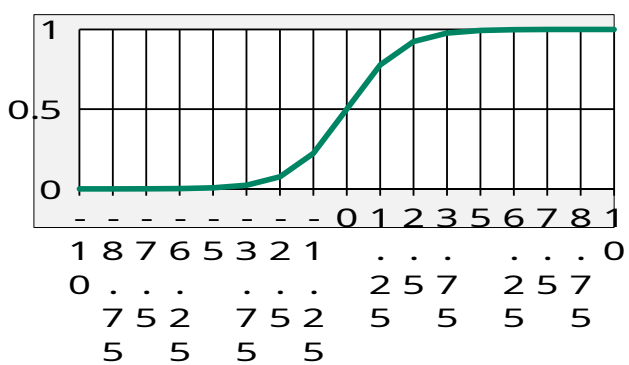

ACTIVATION FUNCTIONS

Linear

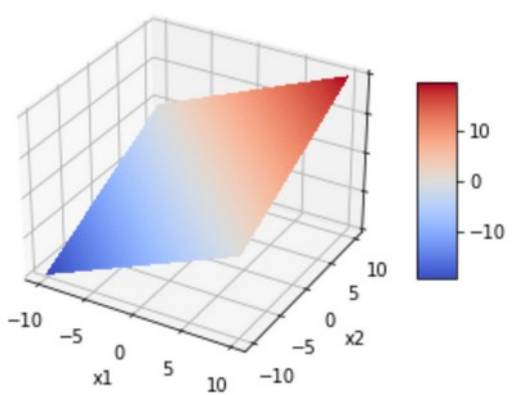

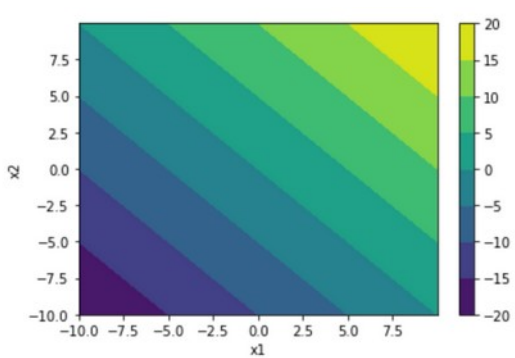

ReLU

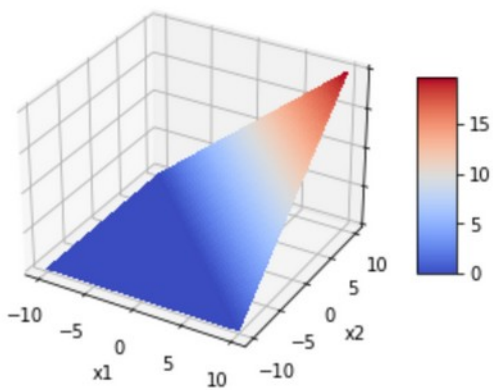

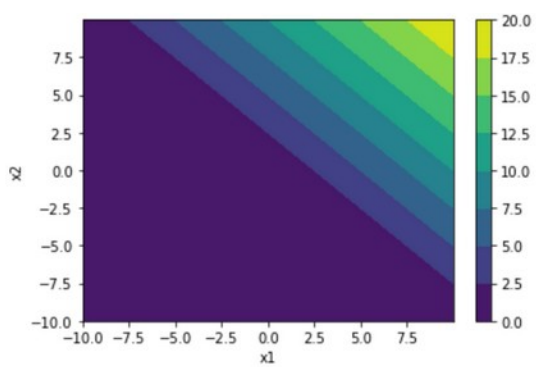

Sigmoid

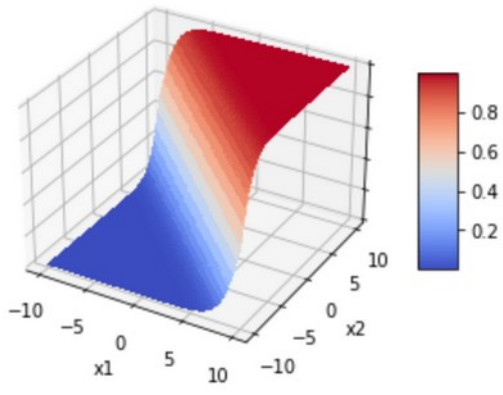

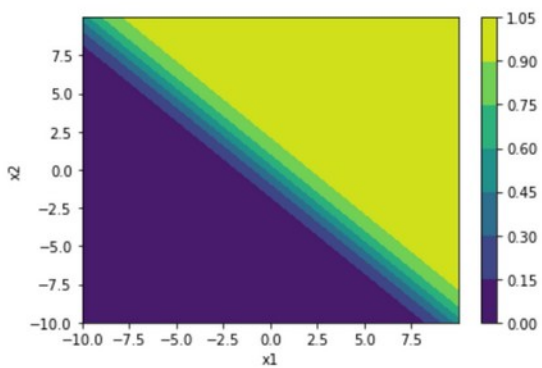

NVIDIA DEEP LEARNING INSTITUTE

ACTIVATION FUNCTIONS

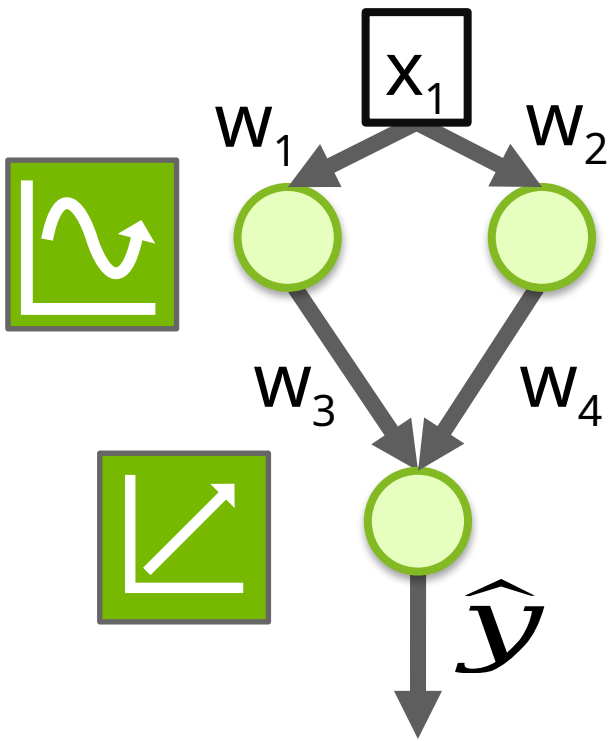

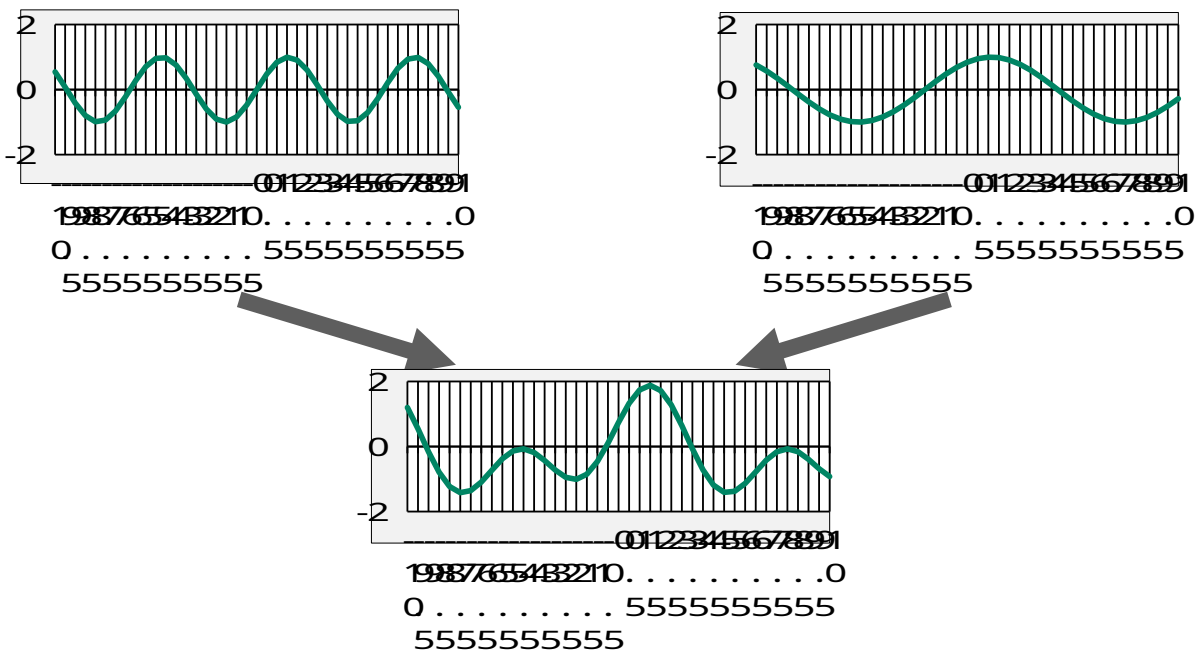

OVERFITTING

OVERFITTING

Why not have a super large neural network?

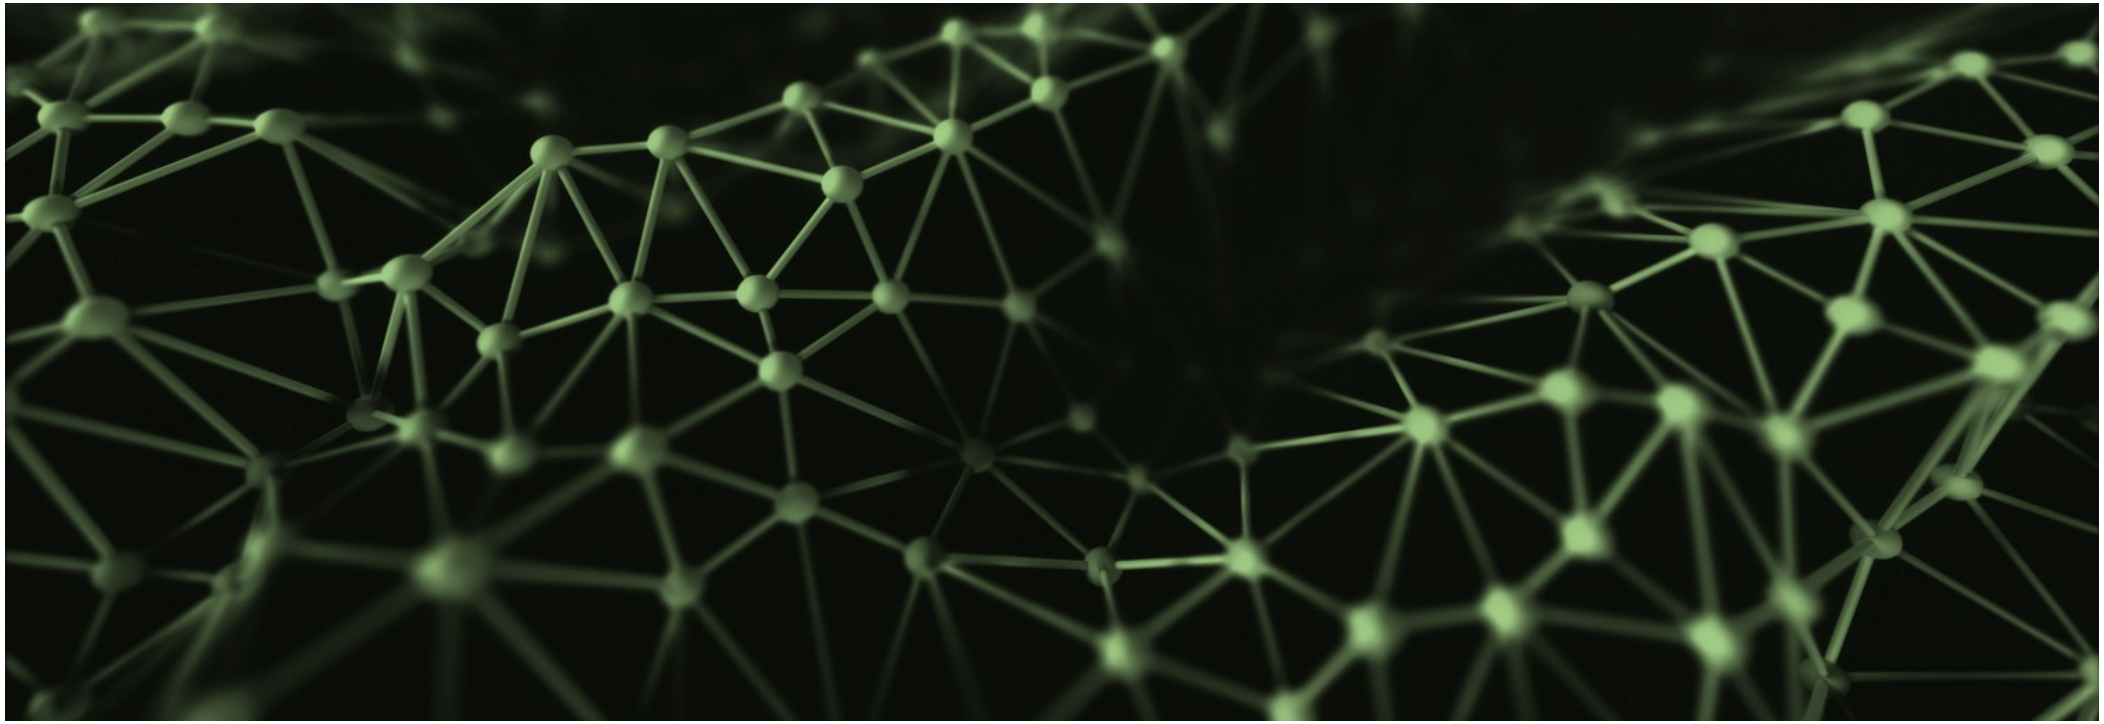

OVERFITTING

Which Trendline is Better?

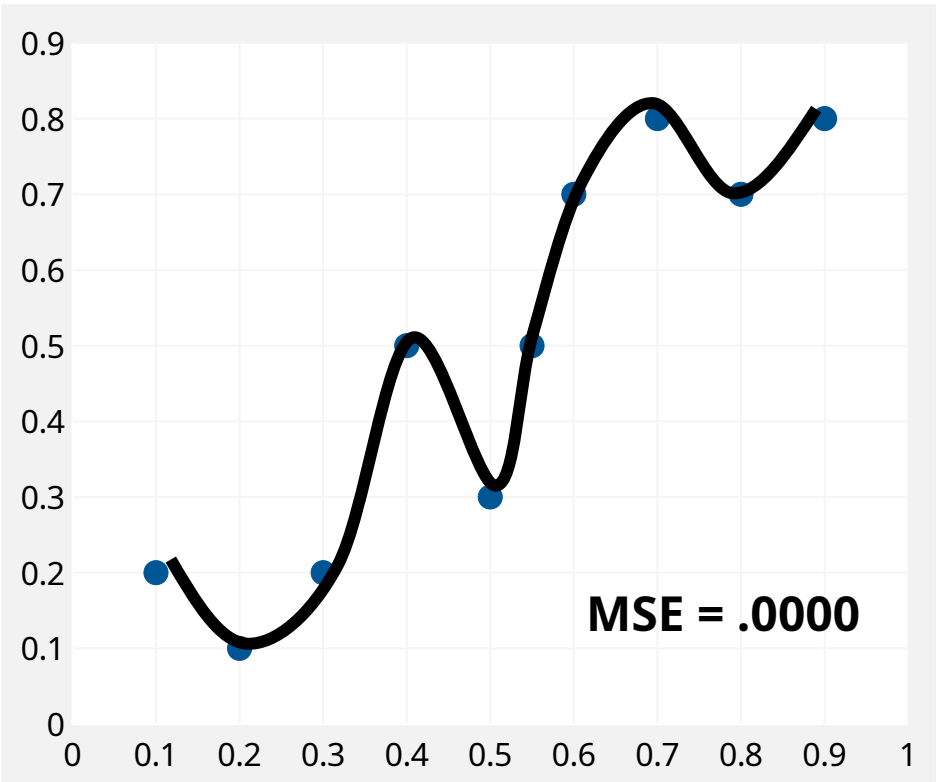

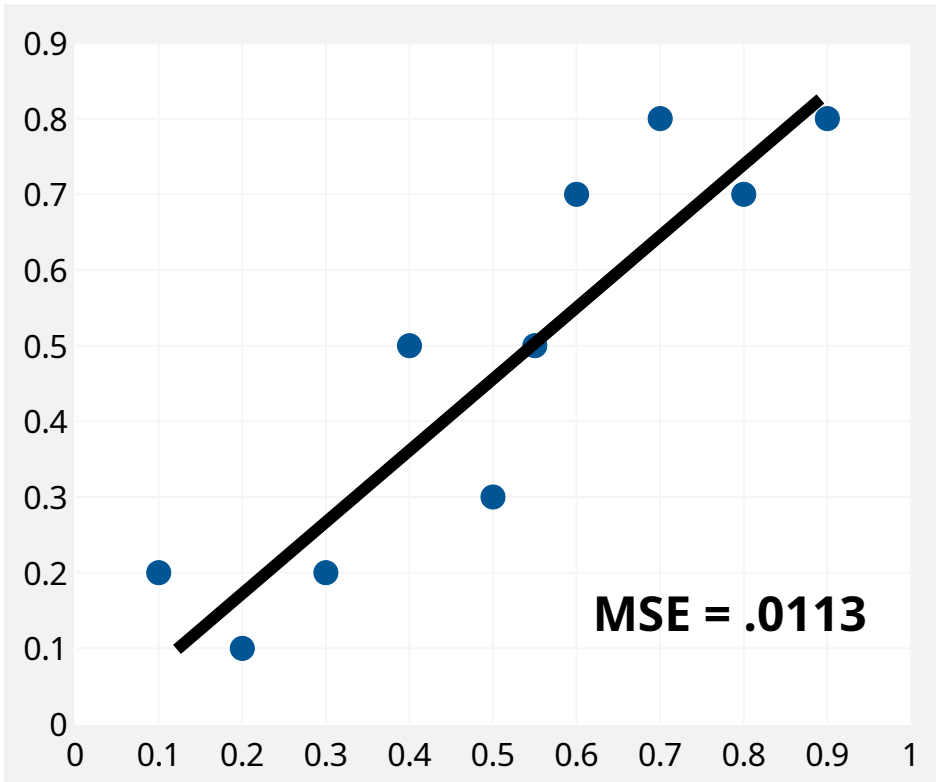

OVERFITTING

Which Trendline is Better?

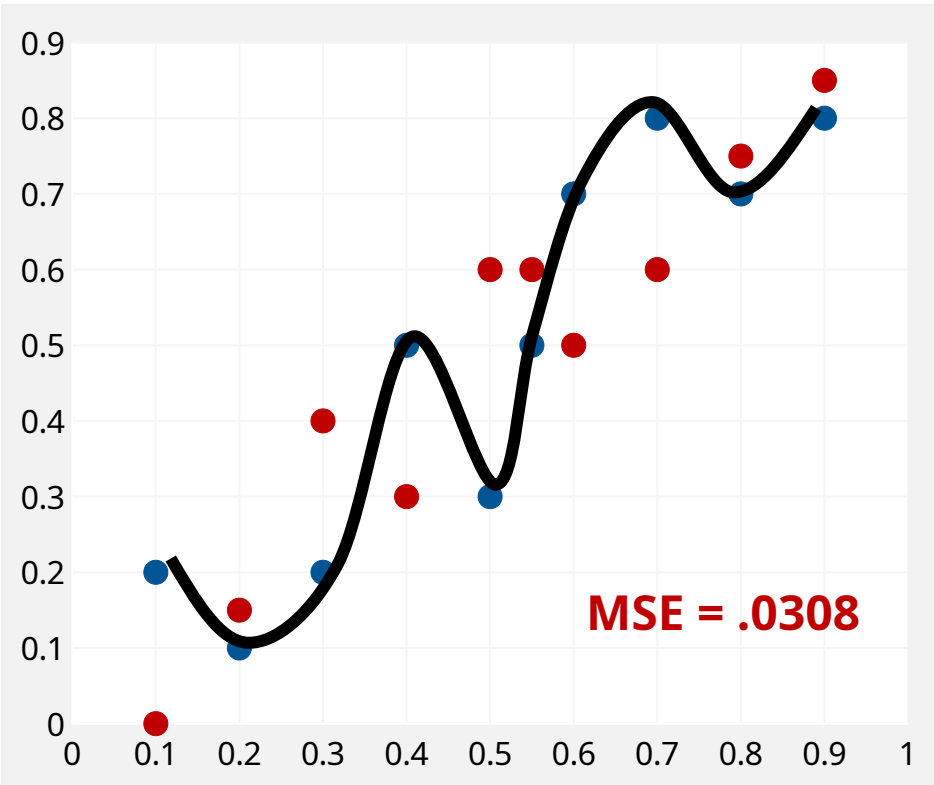

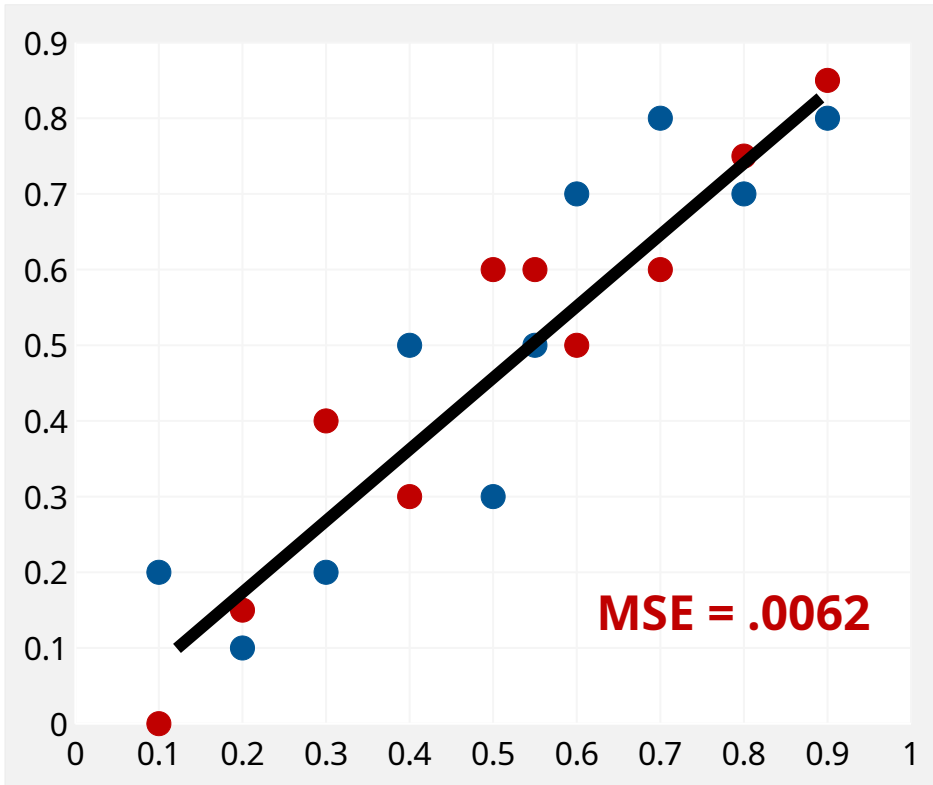

TRAINING VS VALIDATION DATA

Avoid memorization

Training data

- Core dataset for the model to learn on

Validation data

- New data for model to see if it truly understands (can generalize)

Overfitting

- When model performs well on the training data, but not the validation data (evidence of memorization)
- Ideally the accuracy and loss should be similar between both datasets

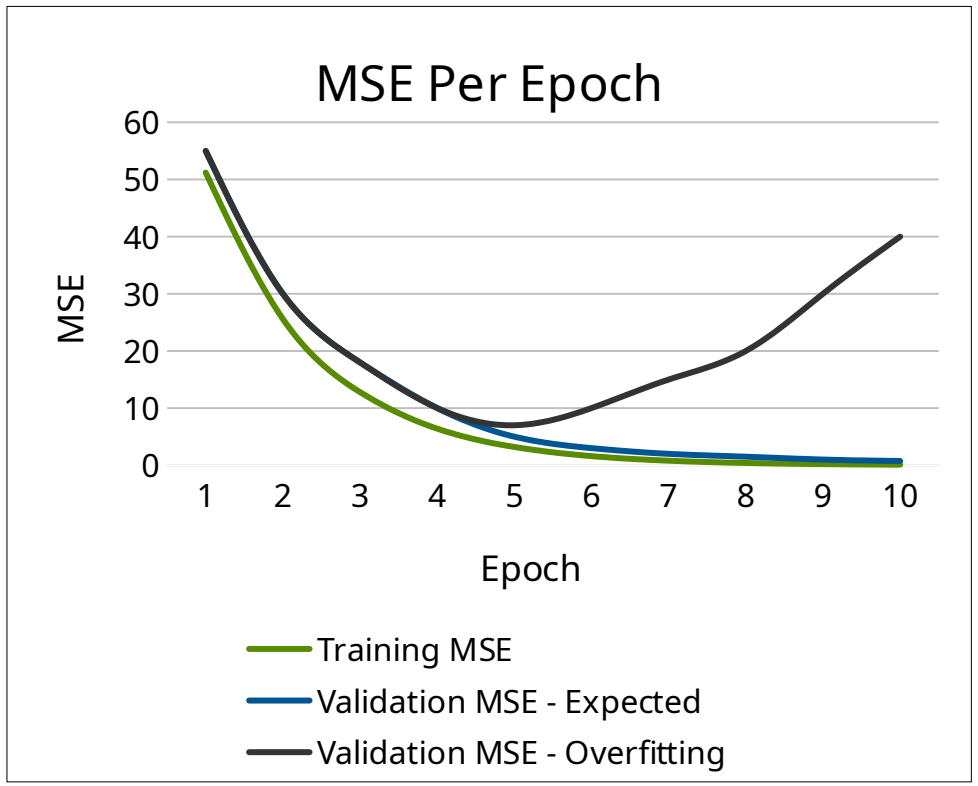

FROM REGRESSION TO CLASSIFICATION

AN MNIST MODEL

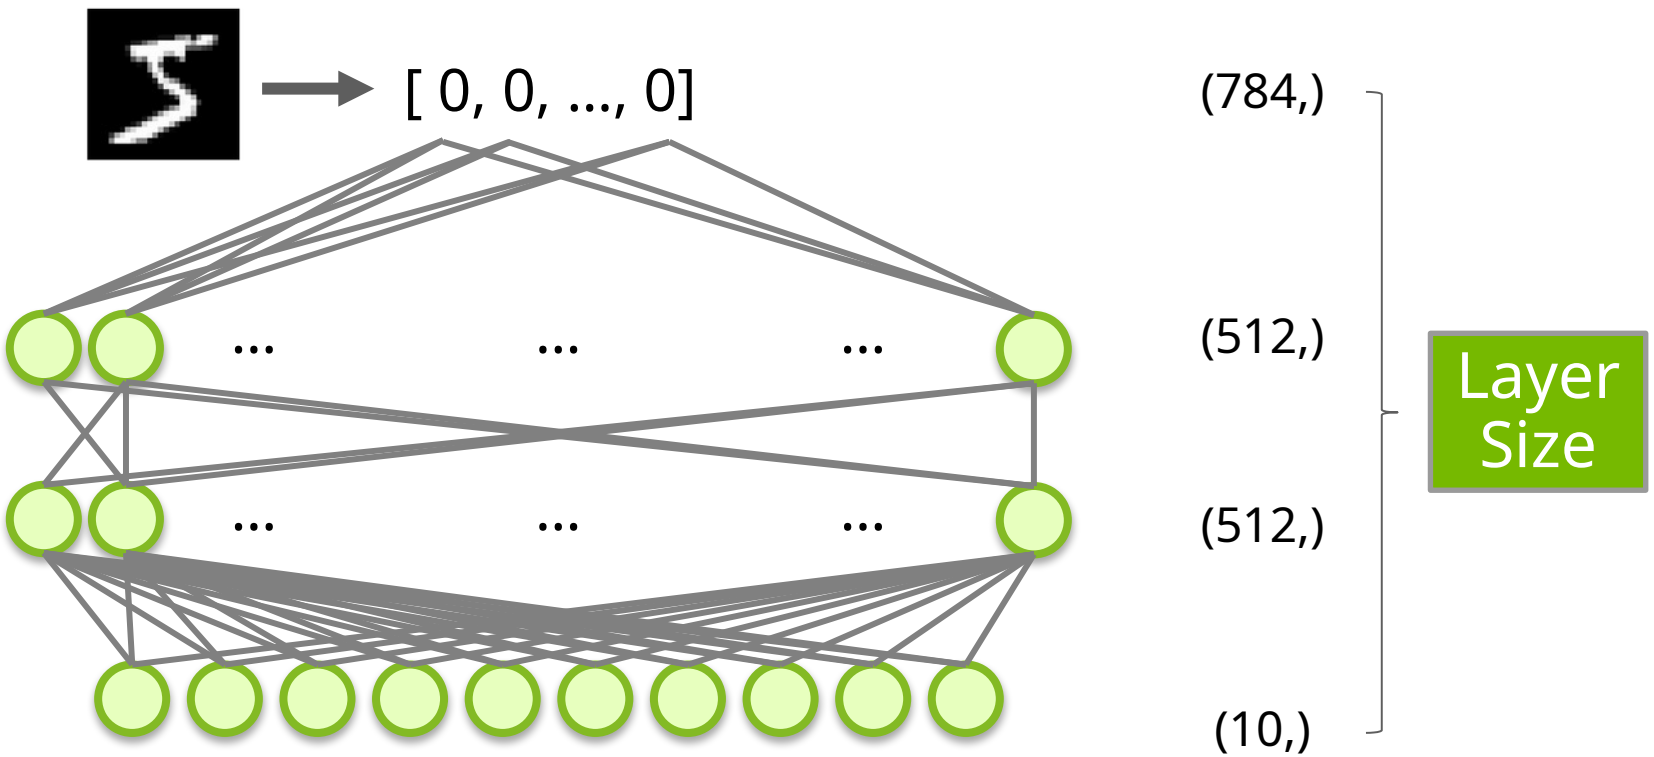

AN MNIST MODEL

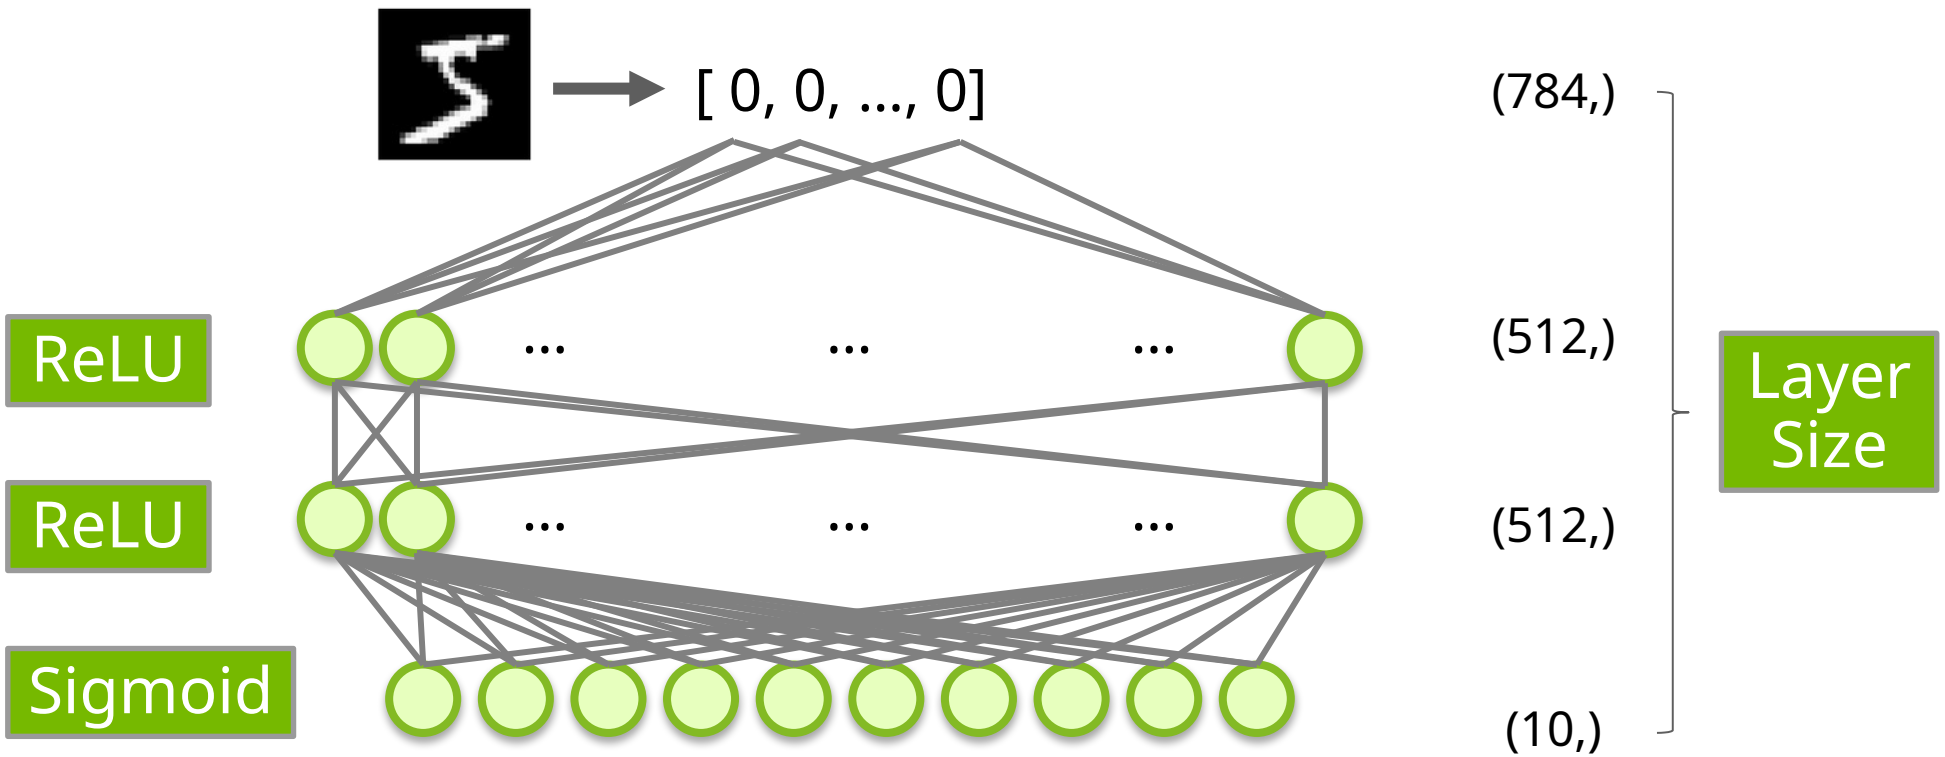

AN MNIST MODEL

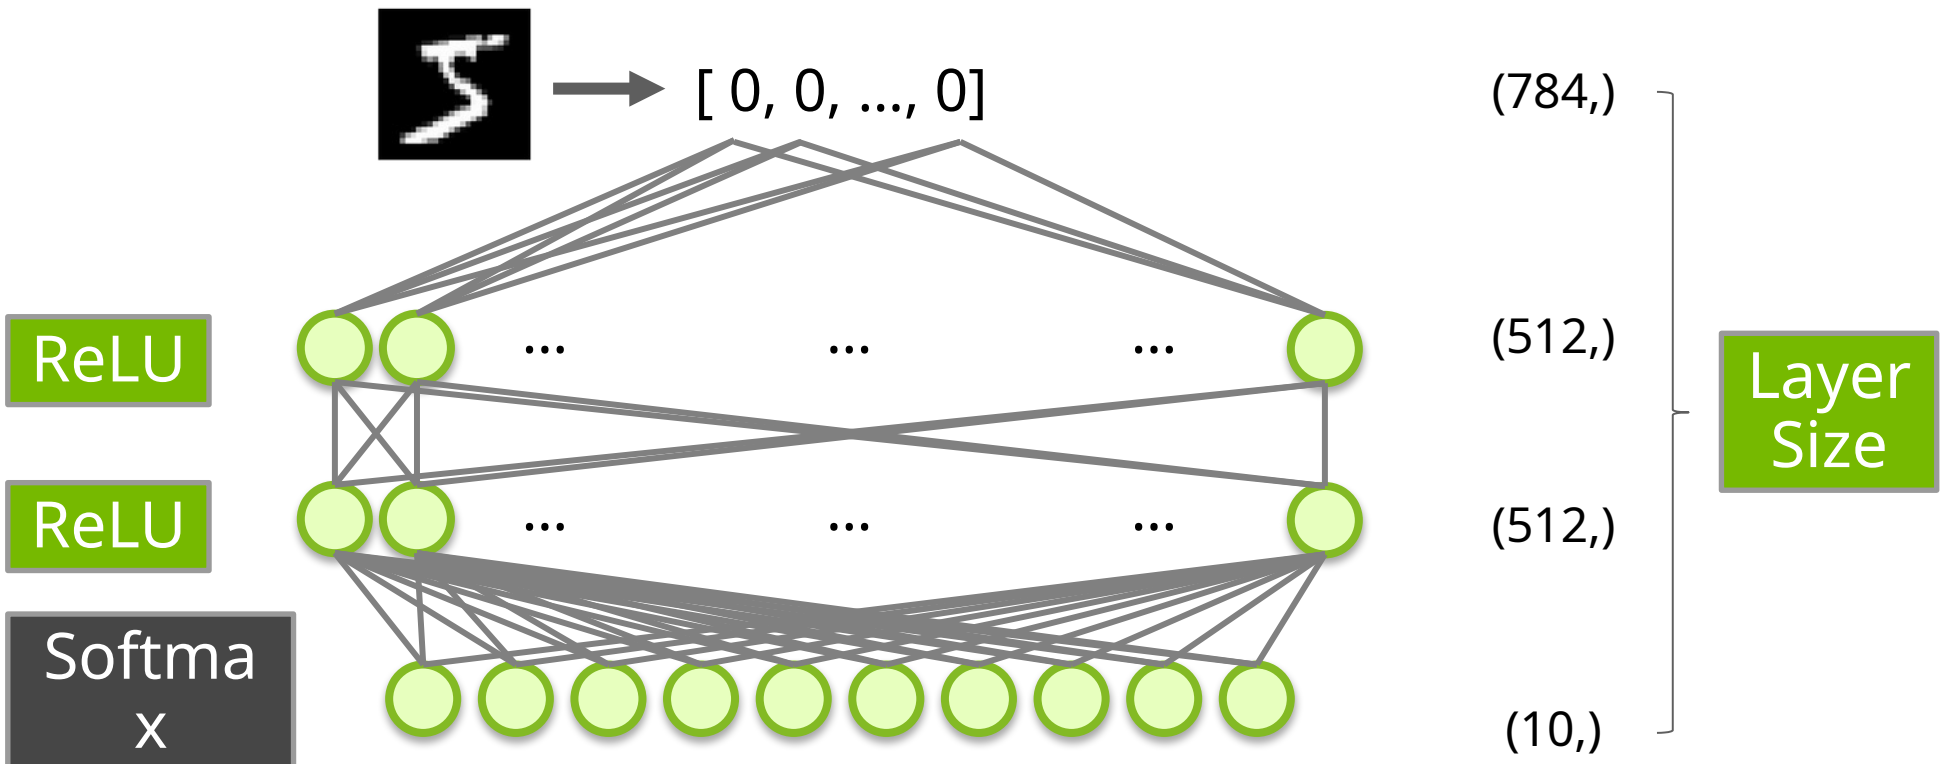

RMSE FOR PROBABILITIES?

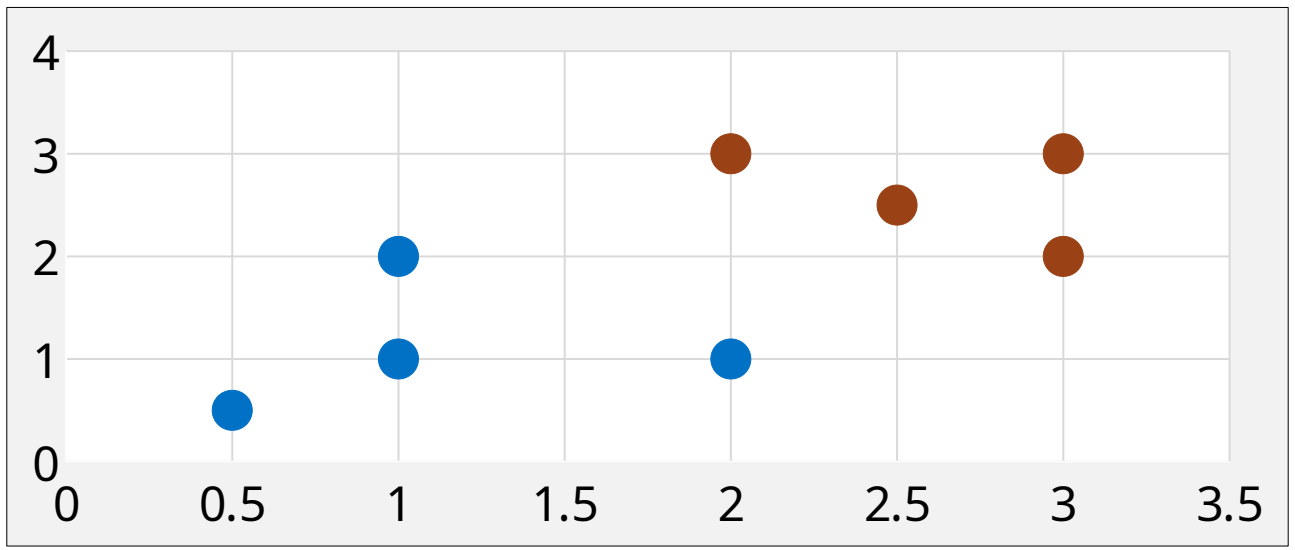

NVIDIA DEEP LEARNING INSTITUTE

RMSE FOR PROBABILITIES?

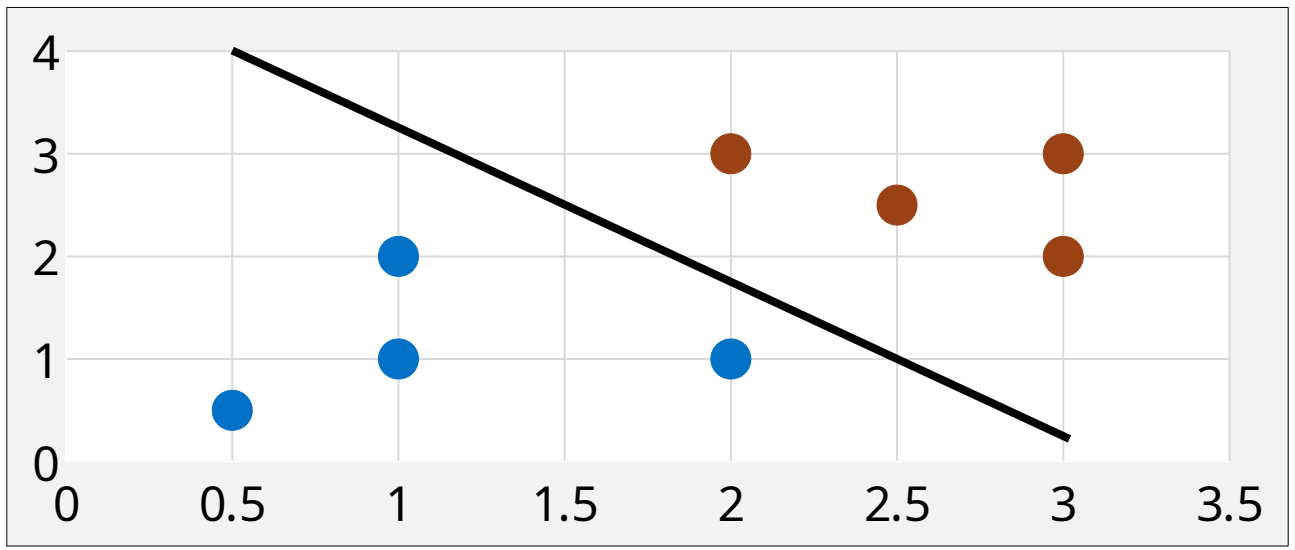

NVIDIA DEEP LEARNING INSTITUTE

CROSS ENTROPY

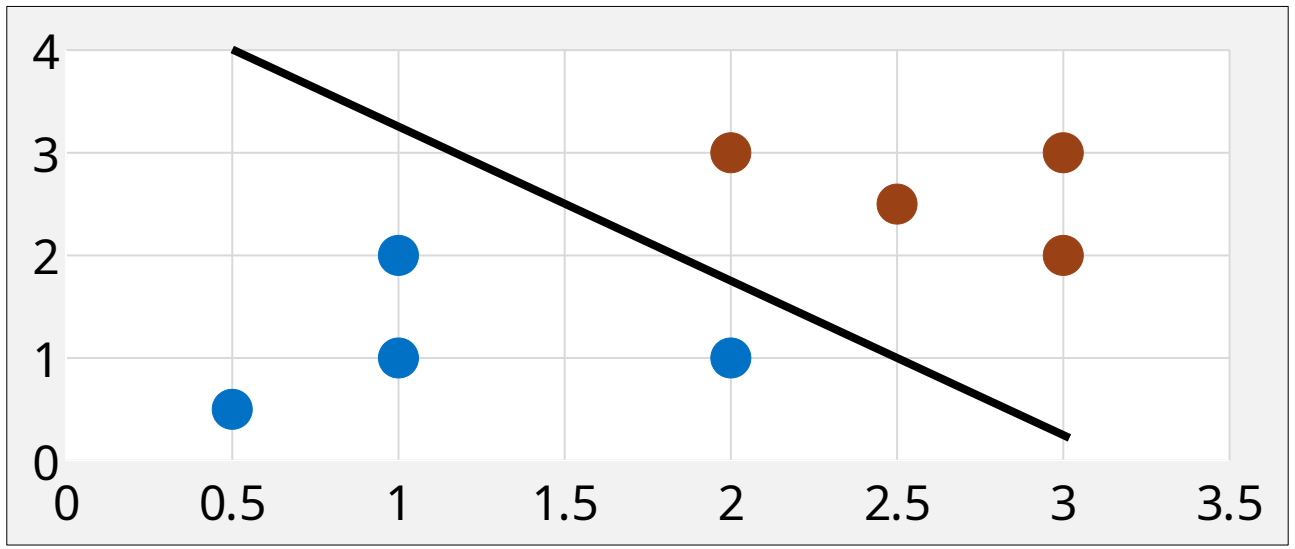

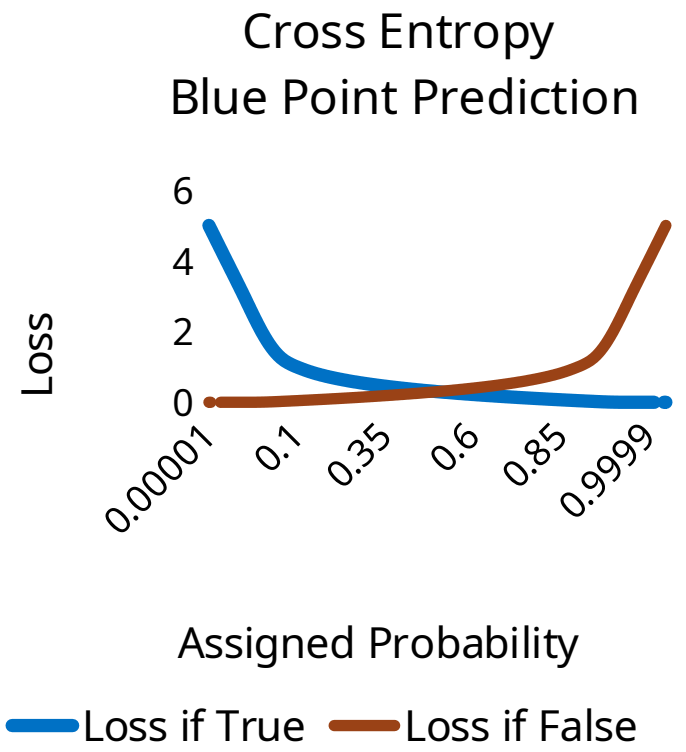

NVIDIA DEEP LEARNING INSTITUTE

CROSS ENTROPY

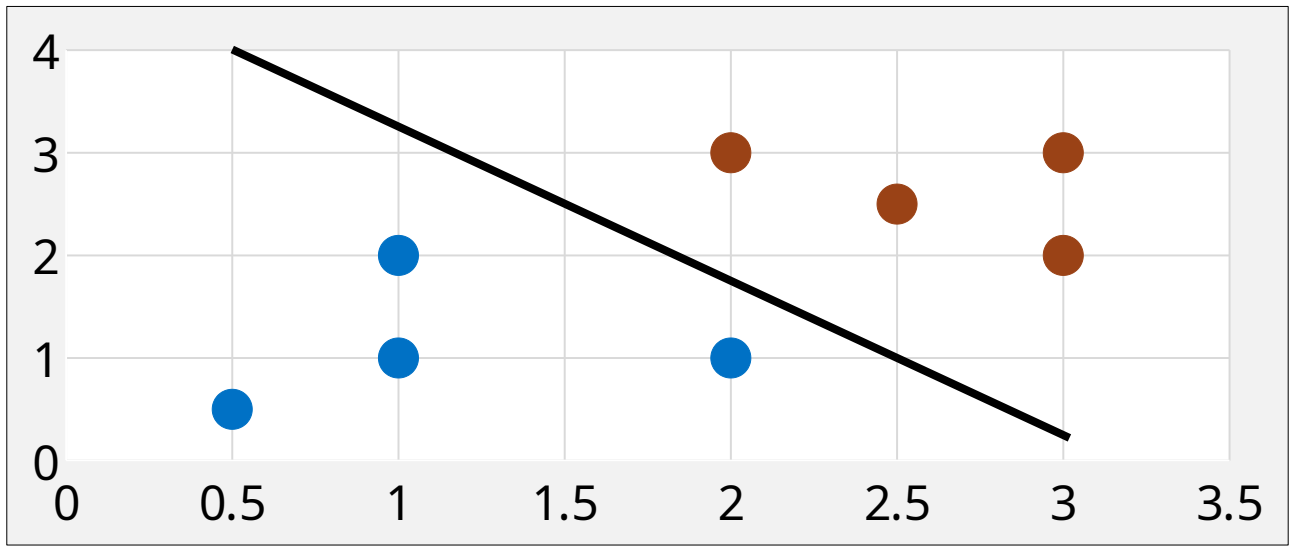

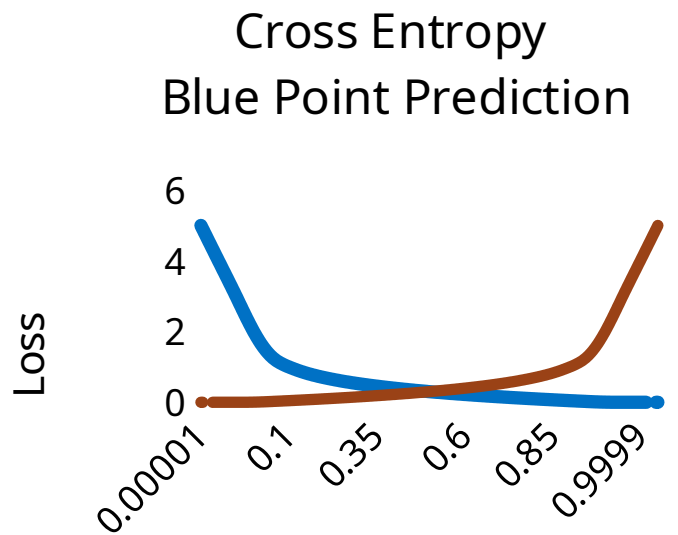

$$Loss = -\left((t(x)\cdot \log\left(p(x)\right) + \left(1 - t(x)\right)\cdot \log\left(1 - p(x)\right)\right)$$

$$t(x) = target\,(0\,if\ False, 1\,if\ True)$$

$$p(x) = probability\ prediction\ of\ point\ x$$

NVIDIA DEEP LEARNING INSTITUTE

CROSS ENTROPY

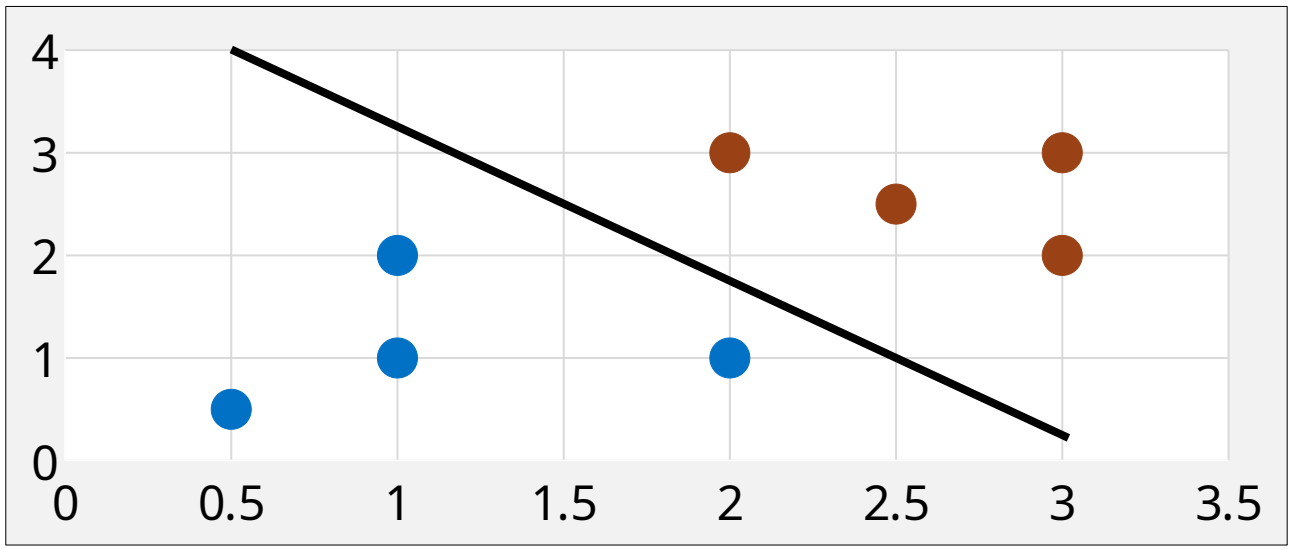


```python
def cross_entropy(y_hat, y_actual):
    """Infinite error for misplaced confidence."""
    loss = log(y_hat) if y_actual else log(1-y_hat)
    return -1*loss
```

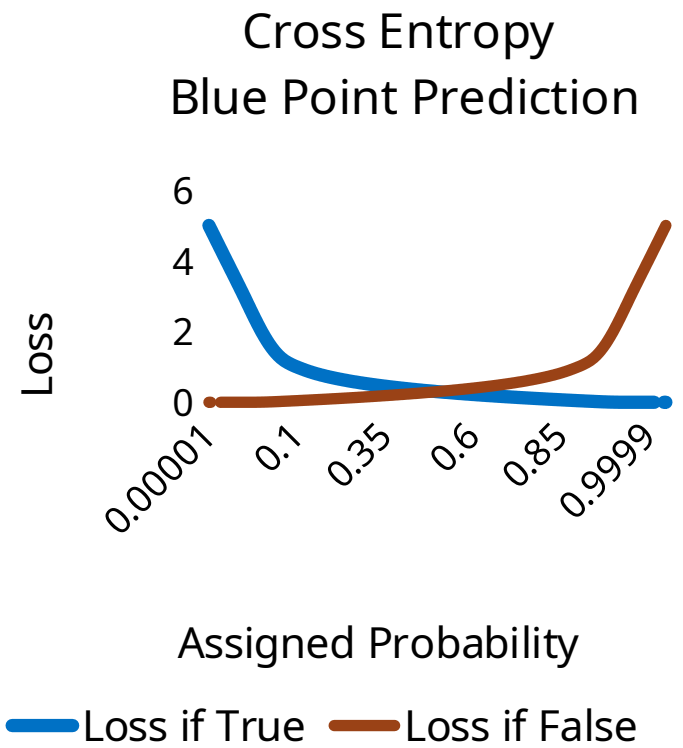

NVIDIA DEEP LEARNING INSTITUTE

BRINGING IT TOGETHER

THE NEXT EXERCISE

The American Sign Language Alphabet

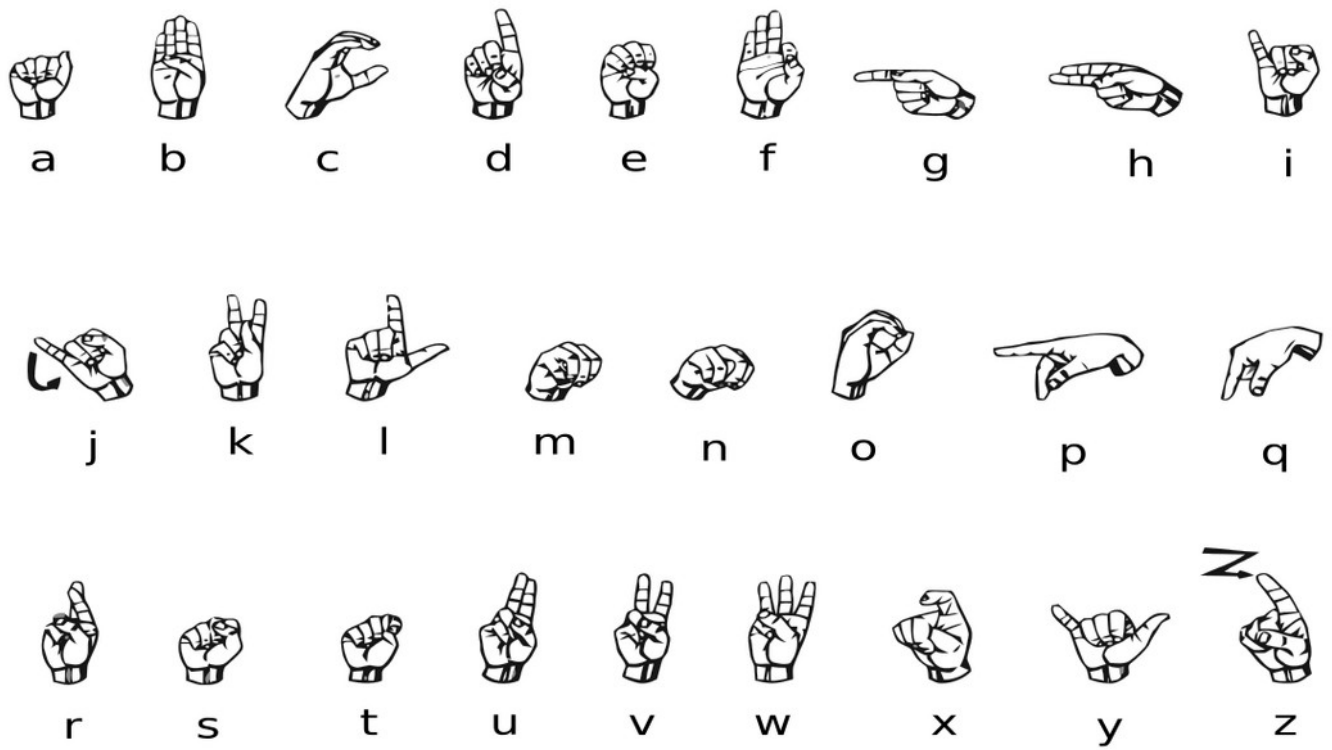

LET’S GO!

NVIDIA NVIDIA DEEP LEARNING INSTITUTE

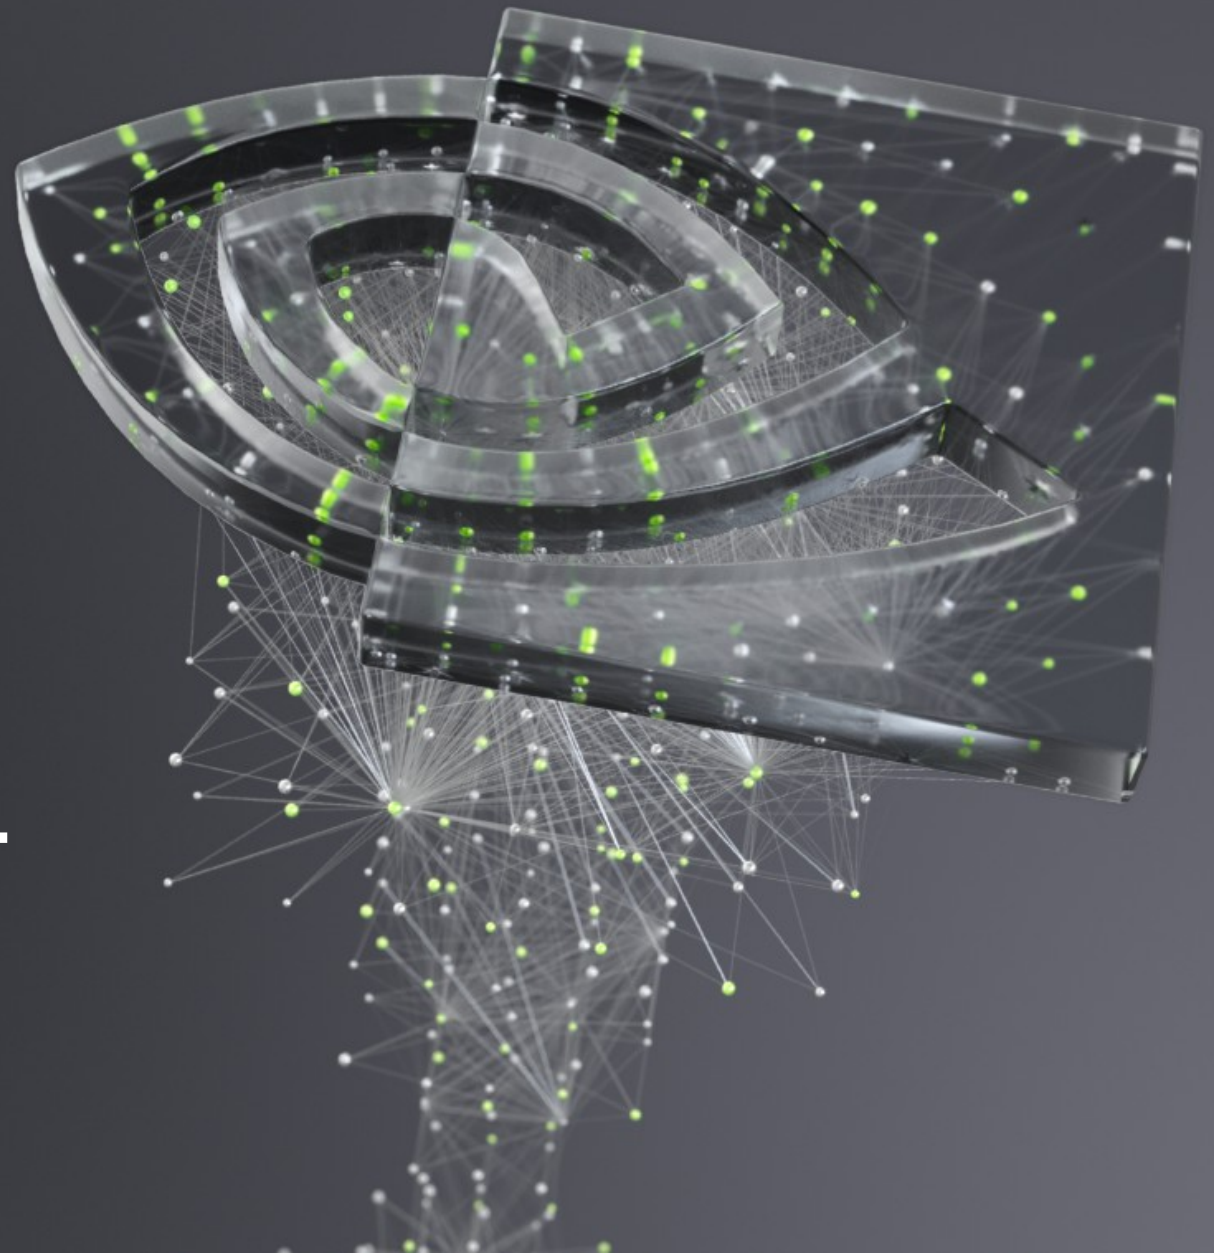

APPENDIX: GRADIENT DESCENT

HELPING THE COMPUTER CHEAT CALCULUS

Learning From Error

$$MSE = \frac{1}{2}((3-(m(1)+b))^2+(5-(m(2)+b))^2)$$

$$\frac{\partial MSE}{\partial m} = 5m+3b-13$$

$$\frac{\partial MSE}{\partial b} = 3m+2b-8$$

$$\frac{\partial MSE}{\partial m} = -3$$

$$\frac{\partial MSE}{\partial b} = -1$$

$$m = -1$$

$$b = 5$$

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

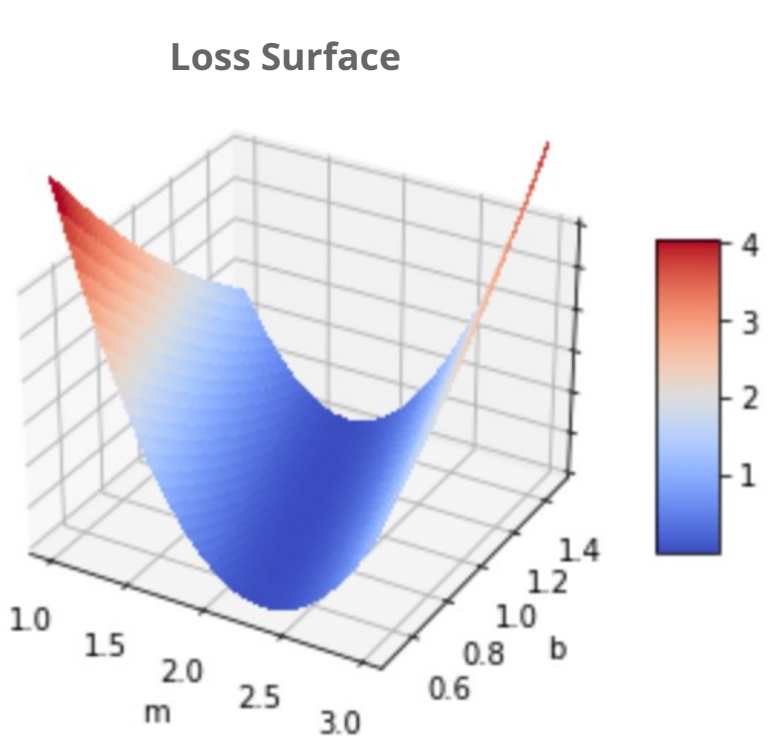

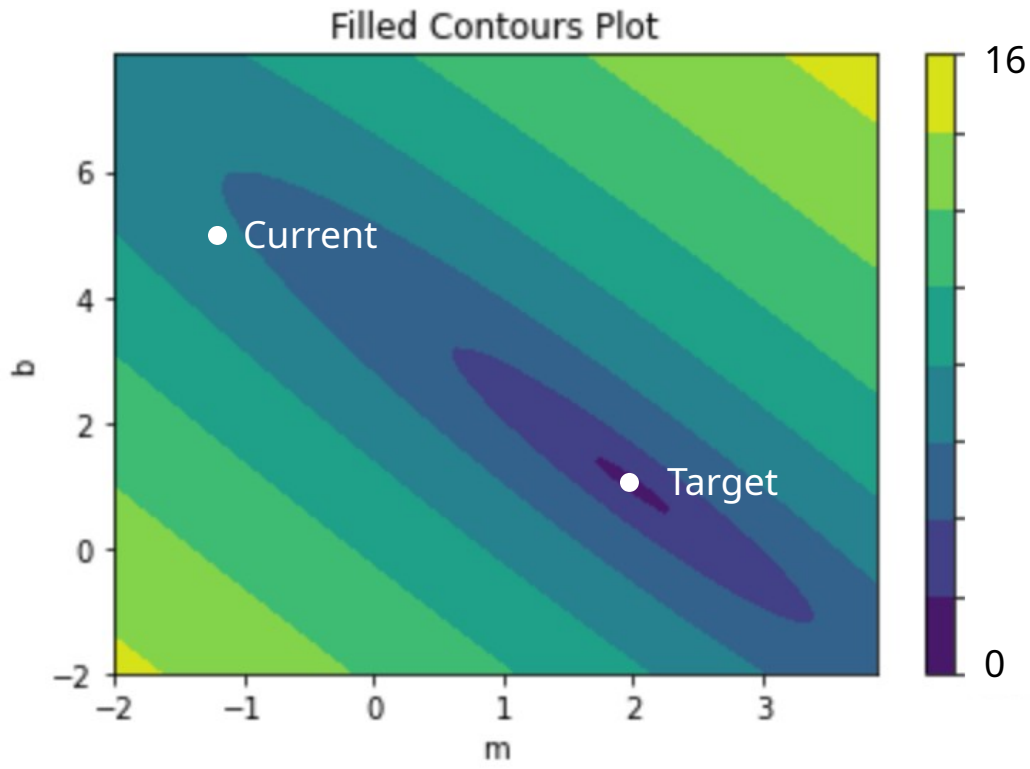

THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \qquad \frac{\partial MSE}{\partial b} = -3$$

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \qquad \frac{\partial MSE}{\partial b} = -3$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$

THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \qquad \frac{\partial MSE}{\partial b} = -3$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$\lambda = .6$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

$$\frac{\partial MSE}{\partial m} = -7 \qquad \frac{\partial MSE}{\partial b} = -3$$

$$m := m - \lambda \frac{\partial MSE}{\partial m}$$

$$\lambda = .005$$

$$b := b - \lambda \frac{\partial MSE}{\partial b}$$

NVIDIA DEEP LEARNING INSTITUTE

THE LOSS CURVE

$$\lambda = .1$$

$$b := 5 + 3\,\lambda = 4.7$$

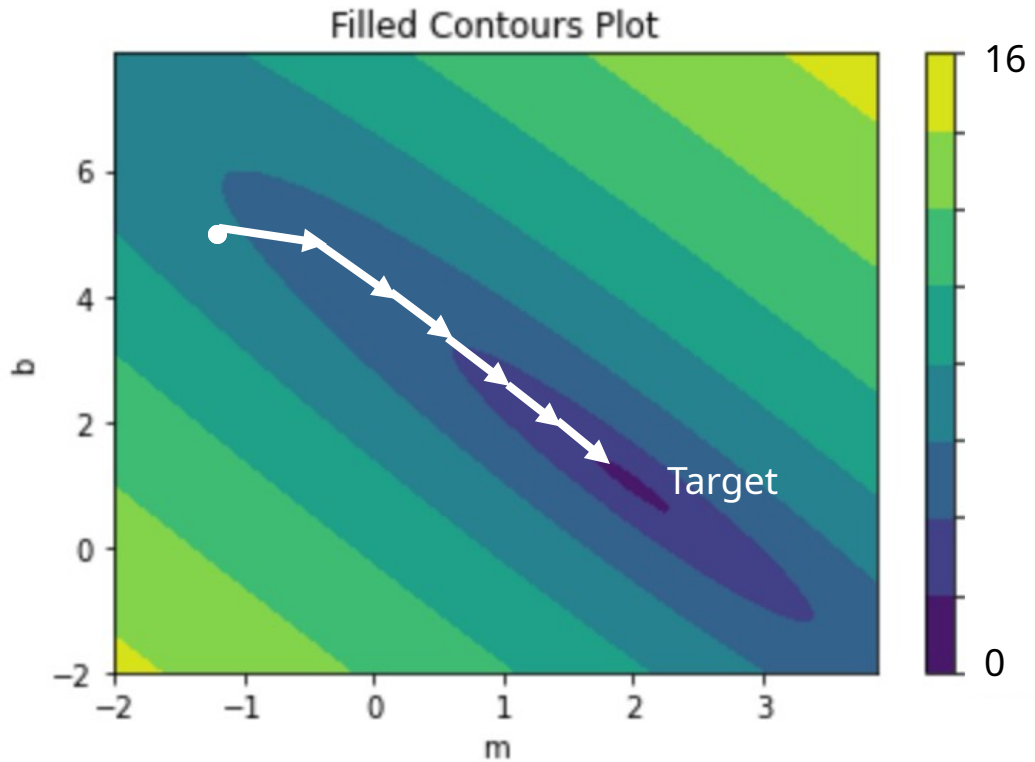

NVIDIA DEEP LEARNING INSTITUTE

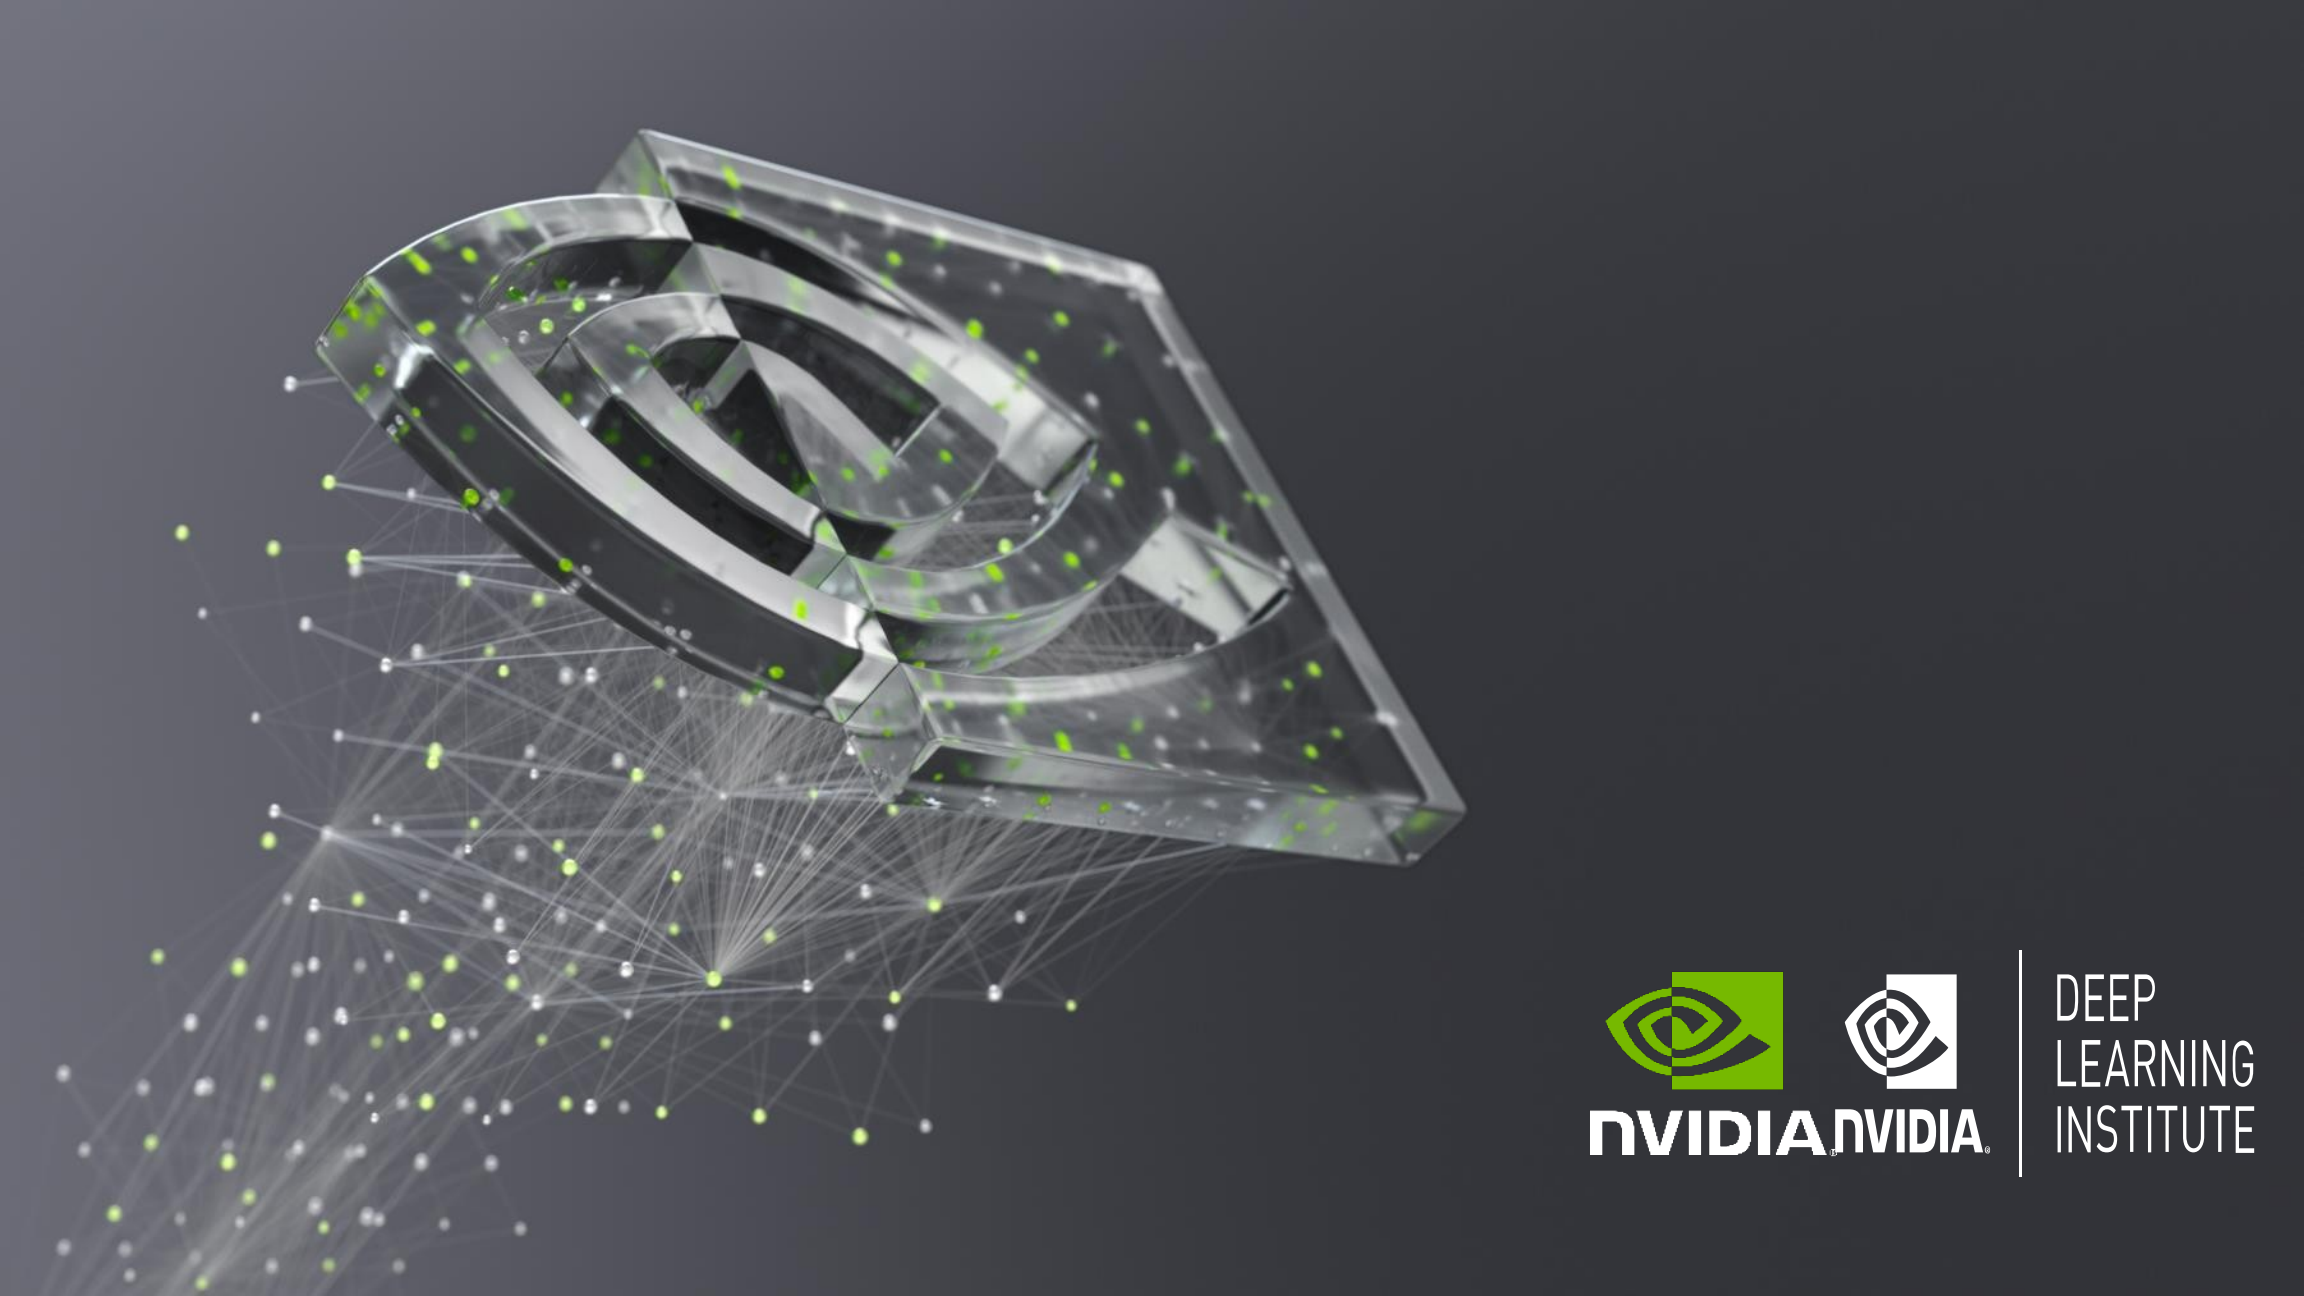
NVIDIA
DEEP
LEARNING
INSTITUTE